

[About](#)[Michael's main blog](#)[Data-driven intelligence blog](#)[Writing](#)

How the Bitcoin protocol actually works

by Michael Nielsen on December 6, 2013

Many thousands of articles have been written purporting to explain Bitcoin peer-to-peer currency. Most of those articles give a hand-wavy account of underlying cryptographic protocol, omitting many details. Even those articles that delve deeper often gloss over crucial points. My aim in this post is to explain the ideas behind the Bitcoin protocol in a clear, easily comprehensible way. We will start from first principles, build up to a broad theoretical understanding of how the protocol works, and then dig down into the nitty-gritty, examining the raw data in a real transaction.

Understanding the protocol in this detailed way is hard work. It is tempting to just take Bitcoin as given, and to engage in speculation about how to get rich with it, whether Bitcoin is a bubble, whether Bitcoin might one day mean the end of fiat currency, and so on. That's fun, but severely limits your understanding. Understanding the details of the Bitcoin protocol opens up otherwise inaccessible vistas. In particular, it is the basis for understanding Bitcoin's built-in scripting language, which makes it possible to use Bitcoin to create new types of financial instruments, such as [contracts](#). New financial instruments can, in turn, be used to create new ways to enable new forms of collective human behaviour. Talk about fun!

I'll describe Bitcoin scripting and concepts such as smart contracts in future posts. This post concentrates on explaining the nuts-and-bolts of the Bitcoin protocol. To understand the post, you need to be comfortable with **public key cryptography** with the closely related idea of **digital signatures**. I'll also assume you're familiar with **cryptographic hashing**. None of this is especially difficult. The basic ideas are taught in freshman university mathematics or computer science classes. They are beautiful, so if you're not familiar with them, I recommend taking a few courses to become familiar.

It may seem surprising that Bitcoin's basis is cryptography. Isn't Bitcoin a currency? A way of sending secret messages? In fact, the problems Bitcoin needs to solve are largely about securing transactions — making sure people can't steal from each other, or impersonate one another, and so on. In the world of atoms we achieve security with devices such as locks, safes, signatures, and bank vaults. In the world of bits we achieve this kind of security with cryptography. And that's why Bitcoin needs a cryptographic protocol.

My strategy in the post is to build Bitcoin up in stages. I'll begin by explaining a simple digital currency, based on ideas that are almost obvious. We'll call this currency *Infocoin*, to distinguish it from Bitcoin. Of course, our first version will have many deficiencies, and so we'll go through several iterations of it. In each iteration introducing just one or two simple new ideas. After several stages of iterations, we'll arrive at the full Bitcoin protocol. We will have reinvented Bitcoin.

This strategy is slower than if I explained the entire Bitcoin protocol in one go. While you can understand the mechanics of Bitcoin through such a one-shot explanation, it would be difficult to understand *why* Bitcoin is designed the way it is. The advantage of the slower iterative explanation is that it gives us a much deeper understanding of each element of Bitcoin.

Finally, I should mention that I'm a relative newcomer to Bitcoin. I've been interested in it loosely since 2011 (and cryptocurrencies since the late 1990s), but only got into the details of the Bitcoin protocol earlier this year. So I'd certainly appreciate corrections of any misapprehensions on my part. Also in the post I've included a number of "problems for the author" – notes to myself about questions that

during the writing. You may find these interesting, but you can also skip them without losing track of the main text.

First steps: a signed letter of intent

So how can we design a digital currency?

On the face of it, a digital currency sounds impossible. Suppose some person call her Alice – has some digital money which she wants to spend. If Alice uses a string of bits as money, how can we prevent her from using the same bit string over and over, thus minting an infinite supply of money? Or, if we can somehow solve this problem, how can we prevent someone else forging such a string of bits, and thus that to steal from Alice?

These are just two of the many problems that must be overcome in order to use digital information as money.

As a first version of Infocoin, let's find a way that Alice can use a string of bits as a (very primitive and incomplete) form of money, in a way that gives her at least some protection against forgery. Suppose Alice wants to give another person, Bob, one infocoin. To do this, Alice writes down the message "I, Alice, am giving Bob one infocoin". She then digitally signs the message using a private cryptographic key. Finally, Alice announces the signed string of bits to the entire world.

(By the way, I'm using capitalized "Infocoin" to refer to the protocol and general concept, and lowercase "infocoin" to refer to specific denominations of the currency. This similar usage is common, though not universal, in the Bitcoin world.)

This isn't terribly impressive as a prototype digital currency! But it does have some virtues. Anyone in the world (including Bob) can use Alice's public key to verify that Alice really was the person who signed the message "I, Alice, am giving Bob one infocoin". No-one else could have created that bit string, and so Alice can't come back around and say "No, I didn't mean to give Bob an infocoin". So the protocol establishes that Alice truly intends to give Bob one infocoin. The same fact that no one else could compose such a signed message – also gives Alice some limited protection from forgery. Of course, *after* Alice has published her message it's possible

people to duplicate the message, so in that sense forgery is possible. But possible from scratch. These two properties – establishment of intent on A and the limited protection from forgery – are genuinely notable features of protocol.

I haven't (quite) said exactly what digital money *is* in this protocol. To make explicit: it's just the message itself, i.e., the string of bits representing the signed message "I, Alice, am giving Bob one infocoin". Later protocols will be in that all our forms of digital money will be just more and more elaborate versions [1].

Using serial numbers to make coins uniquely identifiable

A problem with the first version of Infocoin is that Alice could keep sending the same signed message over and over. Suppose Bob receives ten copies of the message "I, Alice, am giving Bob one infocoin". Does that mean Alice sent *different* infocoins? Was her message accidentally duplicated? Perhaps she is trying to trick Bob into believing that she had given him ten different infocoins, but the message only proves to the world that she intends to transfer one infocoin.

What we'd like is a way of making infocoins unique. They need a label or a serial number. Alice would sign the message "I, Alice, am giving Bob one infocoin with serial number 8740348". Then, later, Alice could sign the message "I, Alice, am giving Bob one infocoin, with serial number 8770431", and Bob (and everyone else) would know that a different infocoin was being transferred.

To make this scheme work we need a trusted source of serial numbers for infocoins. One way to create such a source is to introduce a *bank*. This bank would provide serial numbers for infocoins, keep track of who has which infocoin, and ensure that transactions really are legitimate,

In more detail, let's suppose Alice goes into the bank, and says "I want to withdraw one infocoin from my account". The bank reduces her account balance by one infocoin, and assigns her a new, never-before used serial number, let's say 1234567. Then, when Alice wants to transfer her infocoin to Bob, she signs the message "I, Alice, am giving Bob one infocoin, with serial number 1234567". But Bob c

accept the infocoin. Instead, he contacts the bank, and verifies that: (a) the coin with that serial number belongs to Alice; and (b) Alice hasn't already spent the infocoin. If both those things are true, then Bob tells the bank he wants to spend the infocoin, and the bank updates their records to show that the infocoin with that serial number is now in Bob's possession, and no longer belongs to Alice.

Making everyone collectively the bank

This last solution looks pretty promising. However, it turns out that we can do something much more ambitious. We can eliminate the bank entirely from the protocol. This changes the nature of the currency considerably. It means there is no longer any single organization in charge of the currency. And when you consider the enormous power a central bank has – control over the money supply – that's a pretty huge change.

The idea is to make it so *everyone* (collectively) is the bank. In particular, we want that everyone using Infocoin keeps a complete record of which infocoins belong to which person. You can think of this as a shared public ledger showing all Infocoin transactions. We'll call this ledger the *block chain*, since that's what the coin record will be called in Bitcoin, once we get to it.

Now, suppose Alice wants to transfer an infocoin to Bob. She signs the message "I, Alice, am giving Bob one infocoin, with serial number 1234567", and gives the message to Bob. Bob can use his copy of the block chain to check that, in fact, the infocoin is Alice's to give. If that checks out then he broadcasts both Alice's message and his acceptance of the transaction to the entire network, and everyone updates their copy of the block chain.

We still have the "where do serial numbers come from" problem, but that turns out to be pretty easy to solve, and so I will defer it to later, in the discussion of Bitcoin. A more challenging problem is that this protocol allows Alice to cheat by double spending the infocoin. She sends the signed message "I, Alice, am giving Bob one infocoin, with serial number 1234567" to Bob, and the message "I, Alice, am giving Charlie one infocoin, with [the same] serial number 1234567" to Charlie. Both Bob and Charlie use their copy of the block chain to verify that the infocoin is Alice's to spend. For

they do this verification at nearly the same time (before they've had a chance to talk to each other), both will find that, yes, the block chain shows the coin belongs to Alice. And so they will both accept the transaction, and also broadcast the acceptance of the transaction. Now there's a problem. How should other people update their block chains? There may be no easy way to achieve a consistent ledger of transactions. And even if everyone can agree on a consistent way to update their block chains, there is still the problem that either Bob or Charlie will be double-spending.

At first glance double spending seems difficult for Alice to pull off. After all, she sends the message first to Bob, then Bob can verify the message, and tell the rest of the network (including Charlie) to update their block chain. Once this happens, Charlie would no longer be fooled by Alice. So there is most likely a brief period of time in which Alice can double spend. However, it's obviously undesirable to have any such a period of time. Worse, there are techniques that could be used to make that period longer. She could, for example, use network analysis to find times when Bob and Charlie are likely to have a lot of later communication. Or perhaps she could do something to deliberately disrupt communications. If she can slow communication even a little that makes her double spending much easier.

How can we address the problem of double spending? The obvious solution is that when Alice sends Bob an infocoin, Bob shouldn't try to verify the transaction immediately. Rather, he should broadcast the possible transaction to the entire network of users, and ask them to help determine whether the transaction is legitimate. If everyone collectively decides that the transaction is okay, then Bob can accept the infocoin and everyone will update their block chain. This type of protocol can help prevent double spending, since if Alice tries to spend her infocoin with both Bob and Charlie, the people on the network will notice, and network users will tell both Bob and Charlie that there is a problem with the transaction, and the transaction shouldn't go through.

In more detail, let's suppose Alice wants to give Bob an infocoin. As before, she sends the message "I, Alice, am giving Bob one infocoin, with serial number 123456789." and gives the signed message to Bob. Also as before, Bob does a sanity check by copying the block chain to check that, indeed, the coin currently belongs to Alice. At that point the protocol is modified. Bob doesn't just go ahead and accept the transaction.

transaction. Instead, he broadcasts Alice's message to the entire network. All members of the network check to see whether Alice owns that infocoin. If so, they broadcast the message "Yes, Alice owns infocoin 1234567, it can now be spent by Bob." Once enough people have broadcast that message, everyone updates the block chain to show that infocoin 1234567 now belongs to Bob, and the transaction is complete.

This protocol has many imprecise elements at present. For instance, what does it mean to say "once enough people have broadcast that message"? What does "enough" mean here? It can't mean everyone in the network, since we don't know who is on the Infocoin network. For the same reason, it can't mean some fraction of users in the network. We won't try to make these ideas precise here. Instead, in the next section I'll point out a serious problem with the approach described. Fixing that problem will at the same time have the pleasant side effect of making the ideas above much more precise.

Proof-of-work

Suppose Alice wants to double spend in the network-based protocol I just described. She could do this by taking over the Infocoin network. Let's suppose she uses an automated system to set up a large number of separate identities, let's say 1000, on the Infocoin network. As before, she tries to double spend the same infocoin, sending it both to Bob and Charlie. But when Bob and Charlie ask the network to validate their respective transactions, Alice's sock puppet identities swamp the network, convincing Bob that they've validated his transaction, and to Charlie that they've validated his transaction, possibly fooling one or both into accepting the transaction.

There's a clever way of avoiding this problem, using an idea known as *proof of work*. The idea is counterintuitive and involves a combination of two ideas: (1) to make it *computationally costly* for network users to validate transactions; and (2) to *reward* them for trying to help validate transactions. The reward is used so that anyone on the network will try to help validate transactions, even though that's now made a computationally costly process. The benefit of making it costly to validate transactions is that validation can no longer be influenced by the number of identities someone controls, but only by the total computational power they

to bear on validation. As we'll see, with some clever design we can make it so that a cheater would need enormous computational resources to cheat, making it impractical.

That's the gist of proof-of-work. But to really understand proof-of-work, we need to go through the details.

Suppose Alice broadcasts to the network the news that "I, Alice, am giving Bob one infocoin, with serial number 1234567".

As other people on the network hear that message, each adds it to a queue of pending transactions that they've been told about, but which haven't yet been approved by the network. For instance, another network user named David might have the following queue of pending transactions:

I, Tom, am giving Sue one infocoin, with serial number 1201174.

I, Sydney, am giving Cynthia one infocoin, with serial number 1295618.

I, Alice, am giving Bob one infocoin, with serial number 1234567.

David checks his copy of the block chain, and can see that each transaction is valid. He would like to help out by broadcasting news of that validity to the entire network.

However, before doing that, as part of the validation protocol David is required to solve a hard computational puzzle – the proof-of-work. Without the solution to this puzzle, the rest of the network won't accept his validation of the transaction.

What puzzle does David need to solve? To explain that, let h be a fixed hash function known by everyone in the network – it's built into the protocol. Bitcoin uses the well-known **SHA-256** hash function, but any cryptographically secure hash function would work. Let's give David's queue of pending transactions a label, I , just so it's got a name we can refer to. Suppose David appends a number x (called the *nonce*) to I and computes the combination. For example, if we use $I = \text{"Hello, world!"}$ (obviously this is not a real set of transactions, just a string used for illustrative purposes) and the nonce $x = 1234567890$ (output is in hexadecimal)


```
h("Hello, world!0") =  
1312af178c253f84028d480a6adc1e25e81caa44c749ec81976192e2ec934c64
```

The puzzle David has to solve – the proof-of-work – is to find a nonce x such that when we append x to I and hash the combination the output hash begins with a certain run of zeroes. The puzzle can be made more or less difficult by varying the number of zeroes required to solve the puzzle. A relatively simple proof-of-work puzzle requires just three or four zeroes at the start of the hash, while a more difficult proof-of-work puzzle might require a much longer run of zeros, say 15 consecutive zeroes. In either case, the above attempt to find a suitable nonce, with $x = 0$, is a failure because the output doesn't begin with any zeroes at all. Trying $x = 1$ doesn't work

```
h("Hello, world!1") =  
e9afc424b79e4f6ab42d99c81156d3a17228d6e1eef4139be78e948a9332a7d8
```

We can keep trying different values for the nonce, $x = 2, 3, \dots$. Finally, at $x = 4250$ we obtain:

```
h("Hello, world!4250") =  
0000c3af42fc31103f1fdc0151fa747ff87349a4714df7cc52ea464e12dcd4e9
```

This nonce gives us a string of four zeroes at the beginning of the output hash. This will be enough to solve a simple proof-of-work puzzle, but not enough to solve a more difficult proof-of-work puzzle.

What makes this puzzle hard to solve is the fact that the output from a cryptographic hash function behaves like a random number: change the input even a tiny bit and the output from the hash function changes completely, in a way that's hard to predict. If we want the output hash value to begin with 10 zeroes, say, then David will have to try, on average, to try $16^{10} \approx 10^{12}$ different values for x before he finds a suitable nonce. That's a pretty challenging task, requiring lots of computational power.

Obviously, it's possible to make this puzzle more or less difficult to solve by requiring more or fewer zeroes in the output from the hash function. In fact, the Bitcoin

gets quite a fine level of control over the difficulty of the puzzle, by using a variation on the proof-of-work puzzle described above. Instead of requiring zeroes, the Bitcoin proof-of-work puzzle requires the hash of a block's header to be lower than or equal to a number known as the **target**. This target is automatically adjusted to ensure that a Bitcoin block takes, on average, about ten minutes to validate.

(In practice there is a sizeable randomness in how long it takes to validate: sometimes a new block is validated in just a minute or two, other times it takes many minutes or even longer. It's straightforward to modify the Bitcoin protocol so that the time to validation is much more sharply peaked around ten minutes. Instead of a single puzzle, we can require that multiple puzzles be solved; with some clever design it is possible to considerably reduce the variance in the time to validate transactions.)

Alright, let's suppose David is lucky and finds a suitable nonce, x . (Celebrate! He'll be rewarded for finding the nonce, as described below). He broadcasts the transactions he's approving to the network, together with the value for x . Other participants in the Infocoin network can verify that x is a valid solution to the proof-of-work puzzle. And they then update their block chains to include the new block with the transactions.

For the proof-of-work idea to have any chance of succeeding, network users need an incentive to help validate transactions. Without such an incentive, they have to expend valuable computational power, merely to help validate other people's transactions. And if network users are not willing to expend that power, the system won't work. The solution to this problem is to reward people who help validate transactions. In particular, suppose we reward whoever successfully validates a block of transactions by crediting them with some infocoins. Provided the infocoin reward is large enough that will give them an incentive to participate in validation.


In the Bitcoin protocol, this validation process is called *mining*. For each block of transactions validated, the successful miner receives a bitcoin reward. Initially it was set to be a 50 bitcoin reward. But for every 210,000 validated blocks (which is about once every four years) the reward halves. This has happened just once, to

so the current reward for mining a block is 25 bitcoins. This halving in the reward continues every four years until the year 2140 CE. At that point, the reward will drop below 10^{-8} bitcoins per block. 10^{-8} bitcoins is actually the minimum Bitcoin, and is known as a *satoshi*. So in 2140 CE the total supply of bitcoins will cease to increase. However, that won't eliminate the incentive to help validate transactions. Bitcoin also makes it possible to set aside some currency in a transaction as a *transaction fee*, which goes to the miner who helps validate the transaction. In the early days of Bitcoin transaction fees were mostly set to zero, but as Bitcoin gained in popularity, transaction fees have gradually risen, and are now an additional incentive on top of the 25 bitcoin reward for mining a block.


You can think of proof-of-work as a competition to approve transactions. Entering the competition costs a little bit of computing power. A miner's chance of winning the competition is (roughly, and with some caveats) equal to the proportion of computing power that they control. So, for instance, if a miner controls one percent of the computing power being used to validate Bitcoin transactions, then they have roughly a one percent chance of winning the competition. So provided a lot of computing power is being brought to bear on the competition, a dishonest miner is likely to have only a relatively small chance to corrupt the validation process, even if they expend a huge amount of computing resources.

Of course, while it's encouraging that a dishonest party has only a relative chance to corrupt the block chain, that's not enough to give us confidence in the currency. In particular, we haven't yet conclusively addressed the issue of double spending.

I'll analyse double spending shortly. Before doing that, I want to fill in a little more detail in the description of Infocoin. We'd ideally like the Infocoin network to depend upon the *order* in which transactions have occurred. If we don't have such an order, then at any given moment it may not be clear who owns which infocoins. To solve this we'll require that new blocks always include a pointer to the last block in the chain, in addition to the list of transactions in the block. (The pointer is a hash of the previous block). So typically the block chain is just a linear chain of blocks of transactions, one after the other, with later blocks each containing a pointer to the immediately prior block:




Occasionally, a fork will appear in the block chain. This can happen, for instance two miners happen to validate a block of transactions near-simultaneously, both broadcast their newly-validated block out to the network, and some people update their block chain one way, and others update their block chain the other way.



This causes exactly the problem we're trying to avoid – it's no longer clear what order transactions have occurred, and it may not be clear who owns which transactions. Fortunately, there's a simple idea that can be used to remove any forks. The idea is this: if a fork occurs, people on the network keep track of both forks. But at any time, miners only work to extend whichever fork is longest in their copy of the chain.

Suppose, for example, that we have a fork in which some miners receive block A first and some miners receive block B first. Those miners who receive block A first will continue mining along that fork, while the others will mine along fork B. Let's suppose that the miners working on fork B are the next to successfully mine a block.



After they receive news that this has happened, the miners working on fork A will notice that fork B is now longer, and will switch to working on that fork. From now on, work on fork A will cease, and everyone will be working on the same chain, and block A can be ignored. Of course, any still-pending transactions will still be pending in the queues of the miners working on fork B, and so all transactions will eventually be validated.

Likewise, it may be that the miners working on fork B are the first to extend it. In that case, work on fork A will quickly cease, and again we have a single chain.


No matter what the outcome, this process ensures that the block chain has a unique time ordering of the blocks. In Bitcoin proper, a transaction is not considered confirmed until: (1) it is part of a block in the longest fork, and (2) at least 5 blocks follow it in the longest fork. In this case we say that the transaction has “6 confirmations”. This gives the network time to come to an agreed-upon chain of the blocks. We’ll also use this strategy for Infocoin.

With the time-ordering now understood, let’s return to think about what happens if a dishonest party tries to double spend. Suppose Alice tries to double spend with Bob and Charlie. One possible approach is for her to try to validate a block containing both transactions. Assuming she has one percent of the computing power, she will occasionally get lucky and validate the block by solving the proof-of-work. Unfortunately for Alice, the double spending will be immediately spotted by people in the Infocoin network and rejected, despite solving the proof-of-work problem. So that’s not something we need to worry about.

A more serious problem occurs if she broadcasts two separate transactions. If she spends the same infocoin with Bob and Charlie, respectively. She might, for example, broadcast one transaction to a subset of the miners, and the other transaction to another set of miners, hoping to get both transactions validated. Fortunately, in this case, as we’ve seen, the network will eventually choose one of these transactions, but not both. So, for instance, Bob’s transaction might be confirmed, in which case Bob can go ahead confidently. Meanwhile, Charlie will see that his transaction has not been confirmed, and so will decline Alice’s offer.

this isn't a problem either. In fact, knowing that this will be the case, there is a reason for Alice to try this in the first place.

An important variant on double spending is if Alice = Bob, i.e., Alice tries to spend a coin with Charlie which she is also "spending" with herself (i.e., giving back to herself). This sounds like it ought to be easy to detect and deal with, but, of course, a network to set up multiple identities associated with the same person or organization, so this possibility needs to be considered. In this case, Alice has to wait until Charlie accepts the infocoin, which happens after the transaction has been confirmed 6 times in the longest chain. She will then attempt to fork the chain before the transaction with Charlie, adding a block which includes a transaction in which she pays herself:



Unfortunately for Alice, it's now very difficult for her to catch up with the longest chain. Other miners won't want to help her out, since they'll be working on the longest chain. And unless Alice is able to solve the proof-of-work at least as fast as every other node in the network combined – roughly, that means controlling more than fifty percent of the computing power – then she will just keep falling further and further behind. she might get lucky. We can, for example, imagine a scenario in which Alice controls one percent of the computing power, but happens to get lucky and finds six blocks in a row, before the rest of the network has found any extra blocks. she might be able to get ahead, and get control of the block chain. But this event will occur with probability $1/100^6 = 10^{-12}$. A more general analysis of these lines shows that Alice's probability of ever catching up is infinitesimal unless she is able to solve proof-of-work puzzles at a rate approaching all other nodes combined.

Of course, this is not a rigorous security analysis showing that Alice cannot double spend. It's merely an informal plausibility argument. The [original paper](#) on Bitcoin did not, in fact, contain a rigorous security analysis, only informal a

along the lines I've presented here. The security community is still analysing and trying to understand possible vulnerabilities. You can see some of this [listed here](#), and I mention a few related problems in the "Problems for the future" below. At this point I think it's fair to say that the jury is still out on how secure Bitcoin is.

The proof-of-work and mining ideas give rise to many questions. How much energy is enough to persuade people to mine? How does the change in supply of Bitcoin affect the Infocoin economy? Will Infocoin mining end up concentrated in the hands of a few, or many? If it's just a few, doesn't that endanger the security of the system? Presumably transaction fees will eventually equilibrate – won't this introduce an unwanted source of friction, and make small transactions less desirable? These are all great questions, but beyond the scope of this post. I may come back to them (in the context of Bitcoin) in a future post. For now, we'll stick to our focus on understanding how the Bitcoin protocol works.

Problems for the future

I don't understand why double spending can't be prevented in a simpler way using [two-phase commit](#). Suppose Alice tries to double spend an infocoin by sending it to both Bob and Charlie. The idea is that Bob and Charlie would each broadcast their respective messages to the Infocoin network, along with a request "I want to accept this?" They'd then wait some period – perhaps ten minutes – for any naysayers who could prove that Alice was trying to double spend. If none are heard (and provided there are no signs of attempts to disrupt the network) they'd then accept the transaction. This protocol needs to be hardened against network attacks, but it seems to me to be the core of a good alternative to proof-of-work. How well does this work? What drawbacks and advantages does it have compared to the full Bitcoin protocol?

Early in the section I mentioned that there is a natural way of reducing the variance in time required to validate a block of transactions. If that variance is reduced too much, then it creates an interesting attack possibility. Suppose Alice tries to fork the chain in such a way that: (a) one fork starts with a block in which Alice pays herself, while the other fork starts with a block in which Alice

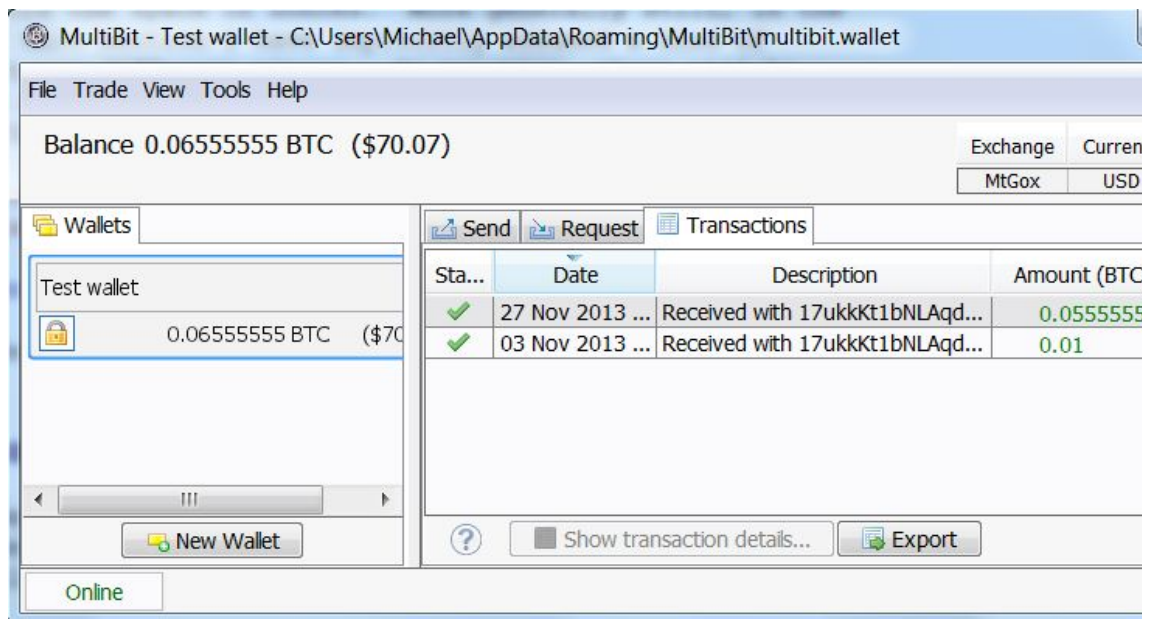
(b) both blocks are announced nearly simultaneously, so roughly half will attempt to mine each fork; (c) Alice uses her mining power to try to mine forks of roughly equal length, mining whichever fork is shorter – this is hard to pull off, but becomes significantly easier if the standard deviation of time-to-validation is much shorter than the network latency; (d) after 5 blocks have been mined on both forks, Alice throws her mining power into the fork more likely that Charles's transaction is confirmed; and (e) after confirming Charles's transaction, she then throws her computational power into the other fork, and attempts to regain the lead. This balancing strategy will have a small chance of success. But while the probability is small, it will certainly be much larger than in the standard protocol, with high variance in the time to validate a block. Is there a way of avoiding this problem?

Suppose Bitcoin mining software always explored nonces starting with $x = 1, x = 2, \dots$. If this is done by all (or even just a substantial fraction) of Bitcoin miners then it creates a vulnerability. Namely, it's possible for miners to improve their odds of solving the proof-of-work merely by starting with a different (much larger) nonce. More generally, it may be possible for attackers to exploit any systematic patterns in the way miners explore the space of nonces. More generally still, in the analysis of this section I have implicitly assumed a symmetry between different miners. In practice, there will be asymmetries, and a thorough security analysis will need to account for those asymmetries.

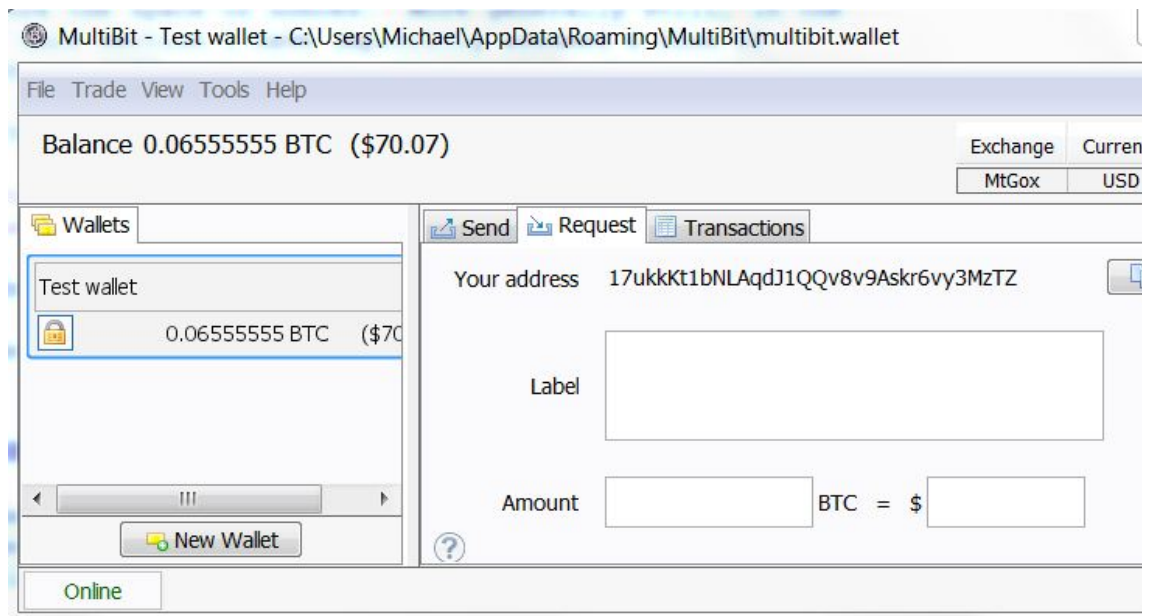
Bitcoin

Let's move away from Infocoin, and describe the actual Bitcoin protocol. There are a few new ideas here, but with one exception (discussed below) they're mostly modifications to Infocoin.

To use Bitcoin in practice, you first install a **wallet** program on your computer. To give you a sense of what that means, here's a screenshot of a wallet called **Mu**. You can see the Bitcoin balance on the left — 0.06555555 Bitcoins, or about 7 dollars at the exchange rate on the day I took this screenshot — and on the right two transactions, which deposited those 0.06555555 Bitcoins:



Suppose you're a merchant who has set up an online store, and you've decided to allow people to pay using Bitcoin. What you do is tell your wallet program a *Bitcoin address*. In response, it will generate a public / private key pair, and hash the public key to form your Bitcoin address:



You then send your Bitcoin address to the person who wants to buy from you. You could do this in email, or even put the address up publicly on a webpage. This is safe since the address is merely a hash of your public key, which can safely be shared with the world anyway. (I'll return later to the question of why the Bitcoin address is a hash and not just the public key.)

The person who is going to pay you then generates a *transaction*. Let's take the data from an **actual transaction** transferring 0.31900000 bitcoins. What below is very nearly the raw data. It's changed in three ways: (1) the data deserialized; (2) line numbers have been added, for ease of reference; and (3) abbreviated various hashes and public keys, just putting in the first six hex digits of each, when in reality they are much longer. Here's the data:

```
1.  {"hash":"7c4025...",
2.   "ver":1,
3.   "vin_sz":1,
4.   "vout_sz":1,
5.   "lock_time":0,
6.   "size":224,
7.   "in":[
8.     {"prev_out":
9.       {"hash":"2007ae...",
10.        "n":0},
11.     "scriptSig":"304502... 042b2d..." } ],
12.  "out":[
13.    {"value":"0.31900000",
14.     "scriptPubKey":"OP_DUP OP_HASH160 a7db6f OP_EQUALVERIFY OP_
```

Let's go through this, line by line.

Line 1 contains the hash of the remainder of the transaction, 7c4025..., expressed in hexadecimal. This is used as an identifier for the transaction.

Line 2 tells us that this is a transaction in version 1 of the Bitcoin protocol.

Lines 3 and 4 tell us that the transaction has one input and one output, respectively. We'll talk below about transactions with more inputs and outputs, and why that's important.

Line 5 contains the value for `lock_time`, which can be used to control when a transaction is finalized. For most Bitcoin transactions being carried out today, this value is 0.

`lock_time` is set to 0, which means the transaction is finalized immediately

Line 6 tells us the size (in bytes) of the transaction. Note that it's not the amount being transferred! That comes later.

Lines 7 through 11 define the input to the transaction. In particular, lines 8 tell us that the input is to be taken from the output from an earlier transaction given `hash`, which is expressed in hexadecimal as `2007ae...`. The `n=0` tells us the first output from that transaction; we'll see soon how multiple outputs (outputs) from a transaction work, so don't worry too much about this for now. Line 9 is the signature of the person sending the money, `304502...`, followed by a space, then the corresponding public key, `04b2d...`. Again, these are both in hexadecimal.

One thing to note about the input is that there's nothing explicitly specifying that the bitcoins from the previous transaction should be spent in this transaction. Instead, the `n=0` tells us that the bitcoins from the `n=0`th output of the previous transaction are spent. So, for example, if the `n=0`th output of the earlier transaction was 2 bitcoins, then 2 bitcoins will be spent in this transaction. This seems like an inconvenient restriction, like trying to buy bread with a 20 dollar note, and not being able to break the note. The solution, of course, is to have a mechanism for providing change. This is done using transactions with multiple inputs and outputs, which we'll discuss in the next section.

Lines 12 through 14 define the output from the transaction. In particular, line 13 tells us the value of the output, 0.319 bitcoins. Line 14 is somewhat complicated; one thing to note is that the string `a7db6f...` is the Bitcoin address of the intended recipient of the funds (written in hexadecimal). In fact, Line 14 is actually a valid expression in Bitcoin's scripting language. I'm not going to describe that language in detail in this post, the important thing to take away now is just that `a7db6f...` is a Bitcoin address.

You can now see, by the way, how Bitcoin addresses the question I swept under the rug in the last section: where do Bitcoin serial numbers come from? In fact, the serial number is played by transaction hashes. In the transaction above, for example, the recipient is receiving 0.319 Bitcoins, which come out of the first

an earlier transaction with hash `2007ae...` (line 9). If you go and look in the chain for that transaction, you'd see that its output comes from a still earlier transaction. And so on.

There are two clever things about using transaction hashes instead of serial numbers. First, in Bitcoin there's not really any separate, persistent "coins" at all, just a series of transactions in the block chain. It's a clever idea to realize that you don't need persistent coins, and can just get by with a ledger of transactions. Second, operating in this way we remove the need for any central authority issuing serial numbers. Instead, the serial numbers can be self-generated, merely by hashing each transaction.

In fact, it's possible to keep following the chain of transactions further back in time. Ultimately, this process must terminate. This can happen in one of two ways. One possibility is that you'll arrive at the very first Bitcoin transaction, contained in a block called **Genesis block**. This is a special transaction, having no inputs, but with one output. In other words, this transaction establishes an initial money supply. The Genesis block is treated separately by Bitcoin clients, and I won't get into that here, although it's along similar lines to the transaction above. You can see the deserialized raw data [here](#), and read about the Genesis block [here](#).

The second possibility when you follow a chain of transactions back in time is that you'll eventually arrive at a so-called *coinbase transaction*. With the exception of the Genesis block, every block of transactions in the block chain starts with a coinbase transaction. This is the transaction rewarding the miner who validated the block of transactions. It uses a similar but not identical format to the transaction above. I won't go through the format in detail, but if you want to see an example, see [here](#). You can read a little more about coinbase transactions [here](#).

Something I haven't been precise about above is what exactly is being signed. The digital signature in line 11. The obvious thing to do is for the payer to sign the transaction (apart from the transaction hash, which, of course, must be generated later). Currently, this is *not* what is done – some pieces of the transaction are signed. This makes some pieces of the transaction **malleable**, i.e., they can be changed. However, this malleability does not include the amounts being paid out, so

recipients, which can't be changed later. I must admit I haven't dug down into details here. I gather that this malleability is under discussion in the Bitcoin community, and there are efforts afoot to reduce or eliminate this malleability.

Transactions with multiple inputs and outputs

In the last section I described how a transaction with a single input and a single output works. In practice, it's often extremely convenient to create Bitcoin transactions with multiple inputs or multiple outputs. I'll talk below about why this can be useful and let's take a look at the data from an **actual transaction**:

```
1. {"hash":"993830...",
2.  "ver":1,
3.  "vin_sz":3,
4.   "vout_sz":2,
5.   "lock_time":0,
6.   "size":552,
7.   "in":[
8.     {"prev_out":{"
9.       "hash":"3beabc...",
10.      "n":0},
11.     "scriptSig":"304402... 04c7d2..."},
12.    {"prev_out":{"
13.      "hash":"fdae9b...",
14.      "n":0},
15.     "scriptSig":"304502... 026e15..."},
16.    {"prev_out":{"
17.      "hash":"20c86b...",
18.      "n":1},
19.     "scriptSig":"304402... 038a52..."}]},
20.  "out":[
21.    {"value":"0.01068000",
22.     "scriptPubKey":"OP_DUP OP_HASH160 e8c306... OP_EQUALVERIFY
```

```

23.      {"value":"4.00000000",
24.      "scriptPubKey":"OP_DUP OP_HASH160 d644e3... OP_EQUALVERIFY

```

Let's go through the data, line by line. It's very similar to the single-input-single-output transaction, so I'll do this pretty quickly.

Line 1 contains the hash of the remainder of the transaction. This is used as a unique identifier for the transaction.

Line 2 tells us that this is a transaction in version 1 of the Bitcoin protocol.

Lines 3 and 4 tell us that the transaction has three inputs and two outputs, respectively.

Line 5 contains the `lock_time`. As in the single-input-single-output case this is 0, which means the transaction is finalized immediately.

Line 6 tells us the size of the transaction in bytes.

Lines 7 through 19 define a list of the inputs to the transaction. Each corresponds to an output from a previous Bitcoin transaction.

The first input is defined in lines 8 through 11.

In particular, lines 8 through 10 tell us that the input is to be taken from the output from the transaction with `hash 3beabc...`. Line 11 contains the signature followed by a space, and then the public key of the person sending the bitcoins.

Lines 12 through 15 define the second input, with a similar format to lines 8 through 11. And lines 16 through 19 define the third input.

Lines 20 through 24 define a list containing the two outputs from the transaction.

The first output is defined in lines 21 and 22. Line 21 tells us the value of the output is 0.01068000 bitcoins. As before, line 22 is an expression in Bitcoin's scripting language. The main thing to take away here is that the string `e8c30622...` is the Bitcoin address of the intended recipient of the funds.

The second output is defined in lines 23 and 24, with a similar format to the first.

One apparent oddity in this description is that although each output has a value associated to it, the inputs do not. Of course, the values of the respective inputs can be found by consulting the corresponding outputs in earlier transactions. In a standard Bitcoin transaction, the sum of all the inputs in the transaction must be at least as much as the sum of all the outputs. (The only exception to this principle is the Genesis block, and in coinbase transactions, both of which add to the overall supply.) If the inputs sum up to more than the outputs, then the excess is called the *transaction fee*. This is paid to whichever miner successfully validates the transaction, making the current transaction a part of a new block.

That's all there is to multiple-input-multiple-output transactions! They're a generalization of single-input-single-output transactions.

One nice application of multiple-input-multiple-output transactions is the idea of *change*. Suppose, for example, that I want to send you 0.15 bitcoins. I can do this by spending money from a previous transaction in which I received 0.2 bitcoins. Of course, I don't want to send you the entire 0.2 bitcoins. The solution is to send you 0.15 bitcoins, and to send 0.05 bitcoins to a Bitcoin address which I own. These 0.05 bitcoins are the change. Of course, it differs a little from the change you might receive in a store, since change in this case is what you pay yourself. But the basic idea is similar.

Conclusion

That completes a basic description of the main ideas behind Bitcoin. Of course, I've omitted many details – this isn't a formal specification. But I have described the main ideas behind the most common use cases for Bitcoin.

While the rules of Bitcoin are simple and easy to understand, that doesn't mean it's easy to understand all the consequences of the rules. There is vastly more that could be said about Bitcoin, and I'll investigate some of these issues in future posts.

For now, though, I'll wrap up by addressing a few loose ends.

How anonymous is Bitcoin? Many people claim that Bitcoin can be used anonymously. This claim has led to the formation of marketplaces such as Silk Road.

(and various successors), which specialize in illegal goods. However, the claim that Bitcoin is anonymous is a myth. The block chain is public, meaning that it's possible for anyone to see every Bitcoin transaction ever. Although Bitcoin addresses are not immediately associated to real-world identities, computer scientists have done a **deal of work** figuring out how to de-anonymize "anonymous" social networks. The block chain is a marvellous target for these techniques. I will be extremely surprised if the great majority of Bitcoin users are not identified with relatively high confidence in the near future. The confidence won't be high enough to achieve complete identification but will be high enough to identify likely targets. Furthermore, identification is retrospective, meaning that someone who bought drugs on Silk Road in 2013 can be identifiable on the basis of the block chain in, say, 2020. These de-anonymization techniques are well known to computer scientists, and, one presumes, the NSA. I would not be at all surprised if the NSA and other agencies have already de-anonymized many users. It is, in fact, ironic that Bitcoin is often touted as anonymous. It's not. Bitcoin is, instead, perhaps the most open and transparent financial system the world has ever seen.

Can you get rich with Bitcoin? Well, maybe. Tim O'Reilly **once said**: "Money is like gas in the car – you need to pay attention or you'll end up on the side of the road. A well-lived life is not a tour of gas stations!" Much of the interest in Bitcoin comes from people whose life mission seems to be to find a *really big* gas station. I admit I find this perplexing. What is, I believe, much more interesting and useful is to think of Bitcoin and other cryptocurrencies as a way of enabling new forms of collective behaviour. That's intellectually fascinating, offers marvellous creative possibilities, is socially valuable, and may just also put some money in the bank. If money in the bank is your primary concern, then I believe that other strategies are much more likely to succeed.

Details I've omitted: Although this post has described the main ideas behind Bitcoin, there are many details I haven't mentioned. One is a nice space-saving trick in the protocol, based on a data structure known as a **Merkle tree**. It's a delightful detail, and worth checking out if fun data structures are your thing. You can get an overview in the **original Bitcoin paper**. Second, I've said little about the **Bitcoin network** – questions like how the network deals with denial of service attacks.

how nodes [join and leave the network](#), and so on. This is a fascinating topic, but it's also something of a mess of details, and so I've omitted it. You can read more about it at some of the links above.

Bitcoin scripting: In this post I've explained Bitcoin as a form of digital, or "programmable", money. But this is only a small part of a much bigger and more interesting story. As we've seen, every Bitcoin transaction is associated to a script in the Bitcoin programming language. The scripts we've seen in this post describe simple transactions like "I gave Bob 10 bitcoins". But the scripting language can also be used to express much more complicated transactions. To put it another way, Bitcoin is *programmable money*. In later posts I will explain the scripting system, and how it is possible to use Bitcoin scripting as a platform to experiment with all sorts of amazing financial instruments.

Thanks for reading. Enjoy the essay? You can tip me with Bitcoin (!) at address [17ukkkKt1bNLAqdJ1QQv8v9Askr6vy3MzTZ](https://blockchain.info/address/17ukkkKt1bNLAqdJ1QQv8v9Askr6vy3MzTZ). You may also enjoy the [first chapter](#) of my forthcoming book on neural networks and deep learning, and may wish to follow [me on Twitter](#).

Footnote

[1] In the United States the question "Is money a form of speech?" is an interesting legal question, because of the protection afforded speech under the US Constitution. In my (legally uninformed) opinion digital money may make this issue more complicated. As we'll see, the Bitcoin protocol is really a way of standing up to the rest of the world (or at least the rest of the Bitcoin network) and avowing "I give such-and-such a number of bitcoins to so-and-so a person" in a way that is extremely difficult to repudiate. At least naively, it looks more like speech than exchanging copper coins, say.

From



Sumedh [permalink](#)

Thanks, I was always too lazy to look up BTC in detail. Your article cleared questions.


I wanted to know one thing what if some smart hacker is able to find some vulnerability in the protocol and he uses that to generate new bitcoins for himself. Once that happens then whole confidence in bitcoins would be gone and it goes to chaos.

Is the above scenario possible?



Bobby [permalink](#)

Your scenario is possible. Just like any other popular piece of software there are incentives for finding exploits, but there are benevolent hackers examining the code to uncover and fix the



Michael Nielsen [permalink](#)

@Bobby: Good point! Yes, that solves much of the problem. The broad point about asymmetries is still true, however. (And is vividly demonstrated by the rise of large mining pools.)

Edit: This is in response to your comment below. I must have used the wrong link when I replied.



achil [permalink](#)


each block starts with a coinbase transaction awarded to the person who solved it. since this transaction is broadcast to all nodes (each node working on the network has

recipient of that transaction), all block in
should'nt (at least not likely) see two bl



gwern [permalink](#)

That bug has actually happened before, but Satoshi/Gavin fixed it before anyone else managed to exploit it. (There have been 2 major Bitcoin bugs that I know of: one allowed you to generate billions of bitcoins, the other allowed you to spend anyone's bitcoins. Neither was fixed before being patched, and there don't seem to have been any others found since.)



Rena [permalink](#)

Interesting. How did those exploits work?



gwern [permalink](#)

I don't know the details of the patches. From the way the exploit worked, it was so bad transactions




Bobby [permalink](#)

I believe have the answer to your third question.

The raw block data that each miner is trying to solve contains a generator transaction. That transaction is where their coins are sent if they solve that block. Because miners competing against each other want their coins to be sent to their own addresses, and those addresses are hashed together with their nonce, it c

matter if everyone starts their nonce from zero. The added randomness from generation transaction addresses prevents each miner from working in the space as others.



Michael [permalink](#)

Thanks Bobby. I had wondered about the same question as the
Your explanation clears it up for me.



JC [permalink](#)

Moreover the nonces need not be enumerable. If randomly picked from a large enough pool it is unlikely that the same nonce gets picked.



Carles [permalink](#)

Very well written!

Only one thing to add (on another post): when you launch Multibit (or bitcoind) where does it connect? a list of IPs? DNS? etc. etc.



Pål Driveklepp [permalink](#)

Bitcoin has 3 methods for finding peers:


<https://en.bitcoin.it/wiki/Network#Bootstrapping>



Benoit [permalink](#)


Thank you for the great write-up!

You write “I’ll return later to the question of why the Bitcoin address is a hash of just the public key”. Did I miss it? Does it have anything to do with quantum computing?



Michael Nielsen [permalink](#)


Oops – actually, I had an extended discussion of this question it just before I posted. The reason I deleted it is that the discussion was inconclusive. The separation seems to be a fairly arbitrary design decision – there are some minor space and security advantages, but not enough (in my opinion) to justify making the Bitcoin address the hash of the public key.



Mark Friedenbach [permalink](#)

There’s a very serious security advantage to reusing addresses until the moment it is spent. That reduces the time a private key could be derived and used to a few minutes. This has significant ramifications for quantum-proof cryptography, if nothing else.

And space-wise we’re talking about saving tens of gigabytes of data from the UTXO set. That’s a pretty significant saving.



Michael Nielsen

Mark,

You’ve described two problems. Both seem like reusing addresses could be a solution. The first problem could be solved by using a single address for all transactions. The second problem could be solved by using a single address for all transactions.

could only ever l
soon; (b) it intro
to publish an ad
when the public
confirmed.


On the second p
savings is less th
the type of trans
other similar sav
seem a bit arbitr

(Actually, it occu
Bitcoin transacti
nice example for



Boldra [permalink](#)

Could the protoc
YAML? It looks l




Benoit [permalink](#)

About your first question:

What would be the incentive for non-miners to answer your question?

Why would you trust the answers or lack thereof?

After all, if I understand correctly, when there is no transaction fee set aside, miners could very well choose to omit transactions from their blocks?




Michael Nielsen [permalink](#)

On incentive: such could be built into the protocol.

On trusting the answers: if someone claims that they see evidence of double spending, you'd require them to present evidence in the signed transaction. The requirement of a signature makes this hard to forge by a malicious naysayer.

On your last point, yes, this is a very interesting question. At present, everything seems to be working okay, but over the long run I suspect the use of Bitcoin for small transactions.



Tom [permalink](#)

On the last point: I could see the transaction taking more time than the time required to confirm a transfer. If it's slower, then you have to pay.

On a related note – what happens if block order?




MG [permalink](#)

I'm curious what happens if/when the underlying crypto is either undermined or broken?

Over the years we've seen flaws that reduce the bits (entropy) in many mechanisms. How would the bitcoin protocol handle, say, a reduction of even 1 bit of difficulty (reduction == 1/2 as hard to attack) .

Also could someone with very large resources overwhelm the network with bad data? Eg, if china wanted to use some super computers or a bot net to stop bitcoin operating by adding all sorts of bad data to the block chains?



Michael Nielsen [permalink](#)

Denial of service type attacks are a real problem. See, e.g: https://en.bitcoin.it/wiki/weaknesses#Denial_of_Service_.28DoS.29

On the first question, the answer is, I think: "That's really complicated and depends on the exact scenario of the break".



cabin [permalink](#)

Android had a bug in their random number api that was successfully exploited. Losing a few bits of entropy won't matter, but in this case they lost nearly all of them.


<http://thegenesisblock.com/security-vulnerability-in-all-android-wallets/>



MG [permalink](#)


One additional question, what happens to bitcoins that are “lost” . ie What the FBI refuses to sign over the bitcoins seized from Silk road, or wallets tl some coins but were lost due to hard drive failure (bad backups) or lost pa Or maybe someone dies but the next of kin doesnt know the details?

The comparison is If I drop \$20 on the ground or my next of kin finds it unc mattress, they can use it.



Michael Nielsen [permalink](#)

Lost bitcoins are just that – gone from the money supply for good. If someone manages to either (a) recover the keypair; or (b) break the underlying crypto.



Mango Cat [permalink](#)

That brings up an interesting scenario, but there has to be some allowance made for re-division of the sub-division of the satoshi.



James H [permalink](#)

With Bitcoin; losing the private key for good is more like accidentally dropping your coins out of an airplane over the pacific ocean.

The private key is crucial to recovering those coins.

Hasitha N. Liyanage [permalink](#)




Looks like we both independently arrived at similar methods of explanation
<http://zen.lk/2013/11/28/how-i-finally-understood-bitcoin/>



Amos [permalink](#)

In the second paragraph of the Bitcoin section, seems it should be 0.0655
0.6555555



Michael Nielsen [permalink](#)

Thanks, fixed.



Vidya [permalink](#)

Thank you so much !!!! I had wanted an understandable primer on Bitcoin
and this was a fabulous read !




Boldra [permalink](#)

I'm surprised that bitcoins in a transaction are a decimal string. It looks like
floating point approximation errors.

Have you read about coin-join? It follows on very nicely from what you've c
here.

I'm looking forward to the next one



Michael Nielsen [permalink](#)

They're not actually a float — as I mention in the article, the m
of Bitcoin is the Satoshi, which is one one hundred millionth of

So it's really specifying an integer number of Satoshis.

(I haven't checked what type is used in the source code; I'd be to know.)



Michael Nielsen [permalink](#)

This page says that it is an integer in the

https://en.bitcoin.it/wiki/Proper_Money_

However, it sounds as though there can be a point / rounding issues with code used on the Bitcoin network.



Uri [permalink](#)


Just wanted to say thanks for a really great essay — the explanation was clear and totally fascinating.



Jim [permalink](#)


Can quantum computers mine bitcoin faster? Does this boil down to how fast a quantum computer can find a string that has a specified property for SHA-256, which we have a quadratic speedup, but probably no more?

I understand that commonly used digital signatures and public-key cryptosystems are broken by quantum computers, so there's not much to be said about that.



Michael Nielsen [permalink](#)

I haven't thought much about it. With that said, I'm pretty sure comments are right – quadratic speedup for finding hash collisions means the asymmetric crypto stuff is broken.



A. Sallai [permalink](#)

This is an incredibly well written article and one that I needed so much. Thank you Michael!



Oli Rhys [permalink](#)

Thanks for this, while I understood the majority of it, the coding element was useful – especially highlighting where the script goes in conjunction with the transaction.

While a lot of people know about bitcoin, there is such a shortage of good quality technical info.

Thanks again 😊




Andrew Jaeger [permalink](#)

Great writeup about how bitcoin functions on a technical level, but I had a question about it as to its use as a currency.

Why is bitcoin built to be inherently deflationary? This seems to be the goal, but against why it will ever gain widespread adoption as a currency. Why does the mining of bitcoin halve every 210,000 blocks? Could there be a point in the future where this is reversed?

Michael Nielsen [permalink](#)



Good questions, and I don't know. I certainly suspect (as do you) that these may ultimately turn out to be design flaws.




Assaf [permalink](#)

Bitcoin is NOT deflationary. It is inflationary with a known and fixed rate up until around 2140 at which point it will stop being inflationary. The only deflation in Bitcoin may happen through coin loss. The same is true for Fiat. The difference is that Fiat can be arbitrarily inflated and with Bitcoin it is not arbitrary.

As far as why inflation is predetermined, this is so Bitcoin is a store of value which is one of the defining properties of a good "money". Why is it inflationary at all (as in, why not start with a predetermined amount of bitcoins that never change). This is part of Bitcoin's decentralized design. Bitcoin designers wanted a way to spread around without starting with a central authority that has them and then throw them out (like, say, ripple). The bitcoin generating part of mining is exactly that.

BTW/ I am yet to see a good argument about why having a monetary system that is a good store of value and does not get diluted by inflation is bad.



KRG [permalink](#)

Inflationary/Deflationary are properties of the supply of real goods. Bitcoin is only a store of value. Real wealth production will gradually slow down until 2140 at the same pace as the drop in Ethereum. Even accounting for the effective tax price of mining, the system needing to increase to attract the necessary capital, a system of directly paying for that service

to be paid out in each transaction to cover actual payments will have to fall to make revenue translates to lower ability to afford

Bitcoin has already seen hyperdeflation but the defined limitations in supply are speculation over it as a commodity that has no predictable future value (never mind use as a measure to reliably store value) rather

What actually needs to be demonstrated is allowing any static, nonproductive accumulation as opposed to using the inherent decline to provide the baseline motivation to use it to store anything beyond cash sufficient to maintain liquidity. There's no justification to use it as a store of value, because value is a property or resource that needs to be accounted for them. Trying to store future production potential is the ultimate fraud and financial manipulation far out of touch with real assets.



mellyra [permalink](#)

Assaf is talking about
talking about price

Both usages are



Jonathan Gold

There are excellent reasons
is the desire to store

Back in 1958 Pa
money as a stor

An Exact Consu
contrivance of r
JPE V66 6 (Dec



Dave [permalink](#)

Actually bitcoin is inherently deflationar
bitcoin economy will grow faster than th
quite intuitive, it does make sense upor
reflects the value of the economy it rep
growing faster than the underlying ecor
money supply is growing slower than th


I think all but a few of us expect the bitc
the supply of bitcoins — hence we have

The wisdom of that choice is another m
many different scenarios for the amoun
new currency entering the system. If bi
then my bet is that these will be signific
among variants.



Wanton [permalink](#)

I'm confused about the block chain. Does everyone have their own versior
they sync to a master? Does every block chain get updated when validatic
completed? Does this mean every person has a record of everyone else's
transactions for ever? Won't this file get really really big?



Michael Nielsen [permalink](#)

In practice, there are thin clients which don't keep a full copy of the chain. But the way the protocol is designed at present there is a small number of people keeping a full copy of the block chain. This is quite a manageable size (about 12 gig). If Bitcoin grows rapidly this may eventually become a problem. There's a nice discussion and related scalability issues here:

<https://en.bitcoin.it/wiki/Scalability>

The conclusion there, which seems to me believable, is that there are many options for scaling Bitcoin at least up to the level at which credit cards are used today, and perhaps further.



Sunny [permalink](#)

Just about the total amount of bitcoins, if I understand well, new bitcoins are generated each time a transaction is processed? It means the more exchanges there are, the more bitcoins in the market there is ?

So the only way to raise the number of bitcoins is to spend some energy on each transaction (that's a little bit wired for me ;-). How were created the first bitcoins there another way of creating bitcoins than checking transactions ?

And thanks a lot for this post because it's really difficult to get a clear picture of it. Regards



Marco [permalink](#)


1. Not per transaction but per block (of transactions).
2. Exchanges are a bad example. The transactions within the exchange happen outside the network. Only if you deposit or withdrawal to/from an exchange, it goes over the network and therefore is

the block chain. There are so many trades going on within an happens internally. And since trades need to happen fast, the not suited for that.

3. yes

4. Google for the 'Genesis Block'. That's how it got started.


5. no



Daniel R. Grayson [permalink](#)

Very nice!

Comments: you use the concept of mining before defining it. Change "pos vulnerabilities" to "possible vulnerabilities". Fix "spending spending money



Michael Nielsen [permalink](#)

Thanks, typos fixed!



Greg [permalink](#)


Excellent article!

I found two typo's which you might want to zap:

"Bob doesn't just go ahead and accept the transaction. Instead, he broadcast message to the entire network." (Change 'broadcast' to 'broadcasts').

"Will Infocoin mining end up in concentrated in the hands of a few, or many? (Remove first 'in').

Michael Nielsen [permalink](#)



Thanks, typos fixed!



Mango Cat [permalink](#)

Thanks for the great Bitcoin writeup. I have gleaned most of what you said from pieces from articles and message boards, it's nice to see it all described in one place.

What I think is more interesting than the cryptography aspect is the social aspect of Bitcoin and why it seems to be succeeding. First big mover and seem to be in-play, and the “anonymous cash” myth also was a big factor, that, I think the carrot of “get paid for solving hard problems (often using of computing resources)” has drawn in many participants who have helped the network by promoting their own self interests.


I balk every time I hear the bit about “every transaction for all users for all time encoded into the block chain” especially when combined with “the chain is secured by solving hard problems”. Scaling this system to support a billion users transacting multiple times per day seems.... unlikely. Your explanation does help to show that problems don't get much harder as transactions scale up – the blocks themselves get larger, but the hash problem doesn't get significantly harder as block size grows unless you start talking about transacting the world's monetary business in a single system, then those blocks would get uncomfortably large in a very short time and the forking problem would be much more complex than choosing between three chains to follow.

I also somewhat disapprove of the concept of encouraging people to “mine for space to earn currency” since that creates an artificial demand for energy that will grow into a significant waste of “real world” resources as such a system scales.

I have been playing “trust network” thought games since the 1980s, I'd like to see a peer-to-peer digital currency system that is based in the concept of trust between identities instead of solving hash problems to get paid. You could still get paid for validating transactions, but it wouldn't have that appeal of “solve more problems to get paid”.

more paid...” plus, a system based on trust would tend to concentrate trust authorities that would be quickly perceived as hopeless to surpass in trust semi-defeating the incentive to participate as a trusted member of the network. Some kind of carrot to the underdogs was included – which would be pure motivational instead of a technically required component.

Anyway, all very interesting to watch. As usual, I got in late and out early (bought around 5, sold around 120, seemed like an awesome profit margin time...) that aspect of Bitcoin is a lot like any other speculative investment certainly fueling interest at this stage.




Michael Nielsen [permalink](#)

On scalability, check out <https://en.bitcoin.it/wiki/Scalability>. There is a lot of useful information there. Like you, though, I wonder about the economics (and impact) of mining.




Michael Newman [permalink](#)

Thanks for writing this great explanation of Bitcoin. I noticed in the first Bitcoin transaction example, you mention 0.39 bitcoins, but the example really describes 0.319 bitcoins, where 0.319 bitcoins goes to one person, and there is a 0.001 bitcoin transaction fee. In other words, did you mean “0.319” instead of “0.39”? At the very least, you need to show “0.31900000” value as an image?




Michael Nielsen [permalink](#)

Thanks, fixed.



typo: “trying to understand possibility vulnerabilities”: possible



Michael Nielsen [permalink](#)

Thanks, fixed.




Jack L [permalink](#)

Thanks for the excellent writeup. I have a question about one item, hopefully I can explain it.

It appears the money you send someone is merely chunks of one or more transactions. Let's say I receive 1 bitcoin at myaddress_123 and I receive 1 bitcoin at myaddress_456. I now want to send you 2 bitcoins from myaddress_789. The previous transactions are the inputs for my transaction to you.

How does the transaction message for the 2 bitcoin transaction prove that I am the recipient of those previous transactions when the addresses are all different? I generate a hash for each input in the new transaction something that can only be generated by whoever was the recipient in the original transactions (myaddress_123 or myaddress_456)?

If the answer is yes, then it seems like unique addresses can be easily linked to a case I don't see any anonymity advantage of using new addresses for each transaction. Sorry if I'm missing something obvious here.



Michael Nielsen [permalink](#)

The proof is in the digital signature. That signature is generated with a private key and a public key which must match (when hashed) the address from the earlier transaction. That proves that the bitcoins are being spent.



ajay [permalink](#)

But (if I understand correctly) the need verified means that you are tied to all y maintain a double life.

If I were a criminal, I might find it very d Stringer, who sells drugs, and Russell f and pillar of the community – and use t bankroll Russell's legitimate businesse with Bitcoin; I can't transfer Stringer's b else in the world knowing about it. Anyo can notice that the flow of money goes Stringer, to Russell.



Jim Lyon [permali](#)

If you really wan would let any cu withdraw value, customers' acco that were the ba originally deposi

Such a bank wo pay interest, ma traditional curren



Space2001 [perm](#)

Extending from ,
more of your bite
transaction. You
transactions eac
one simple scen
into which all the
xxx address as t
necessary amou
set up specifical
middle-man prov
all the incoming
separate out sub
yyyy addresses
association. Mai
on what I have r
provider are unc





Con Kolivas [permalink](#)

Nonce starting at zero is not a vulnerability. Shares are stochastically distributed throughout the 2^{32} nonce range and it makes no difference where you start. The nonce is simply 32 bits out of the whole 320 bit coinbase that you are hashing. There is no way to design a target solution to be distributed anywhere within the range of those 32 bits. If you start at a higher nonce value you simply will have fewer possible chances at finding a solution before you'll need to get/create a new nonce to hash.



Alex F [permalink](#)

I think maybe you're misinterpreting his concern: the danger is that someone can solve blocks *strictly* faster by starting at a different nonce because as you say, the correct nonce could be anywhere in the range $0..2^{32}$ so every guess has the same $1/2^{32}$ chance of being correct.

Instead, the danger is that someone could solve blocks faster than everyone else* if they start at a higher nonce and everyone else starts at 0. Specifically, assuming (on average) everyone can calculate hashes at the same rate, then any transaction whose correct nonce is higher than Y will always be solved first by someone who started at Y, if $0 < Y < 2^{32}$.


Of course this creates an obvious incentive for all participants to guess nonces in a different order than everyone else. So it seems reasonable that most client software would use a random sequence of nonce guesses rather than guessing sequentially from 0. But if we were to find a vulnerability in the random number generator of a client, then it might be possible to design a competing client which, in practice, almost always finds the correct nonce before the target client, by virtue of guessing the same sequence a few steps ahead. This would allow the attacker to successfully validate a share of blocks.

than their actual portion of the collective computational power, of everyone using the vulnerable client and finding the nonce than they should on average.



Gergely Imreh [permalink](#)

I think there's also a "time" field in the p
updated every few seconds. Thus in pr
everyone has the same message and t
everyone has a different message, regi
https://en.bitcoin.it/wiki/Block_hashing_




Michael Nielsen [permalink](#)

Alex has explained my concern well. H
have pointed out (including Gergely) in
differences in the blocks being hashed
sufficient to make this a non-issue.



Rubberman [permalink](#)

A most excellent and well written article! I look forward to more! Thanks.



Bram Stolk [permalink](#)

As people make transactions, the public ledger grows.
Will it not grow to an unmanageable size at some time?


Julian [permalink](#)



I have a couple questions, possibly a subject for a future article.

1. If the block chain forks, do the miners on both sides of the fork keep the block? If so, doesn't it allow someone to continue executing the proof of work even if they know that someone else has solved the proof of work?

2. I am puzzled by transactions in blocks. Is it not possible for two miners to be working on different blocks which contain mostly, although not all, the same transactions? Then, the first one to solve the proof of work will have validated the transactions in the second miner's block. Does the second miner restart his unverified transactions and putting them in a new block?



Michael Nielsen [permalink](#)

On 2, yes: if you're mining and someone else validates some of the transactions you are working on, then you remove them from your block but continue working with the unvalidated transactions.

On 1, it's true that in different forks, different miners will have blocks rewarded. However, over time only one of the forks will become the accepted consensus for confirmed transactions. And so only transactions from one fork will be able to redeem their transactions.




Rob [permalink](#)

Fascinating read, thank you!

One thing I'm having difficulty with is block chain integrity. What will happen if a wallet owner loses his wallet and restores a backup from a few weeks back. He may have some coins, and he may have received some. Those transactions are not in the block chain. How would the block chain get back in sync?

On your question-to-yourself about using two phase commit, I think the major vulnerability would be vulnerability to denial-of-service attack. A malicious user could send

swarm of identities to act as nay-sayers and therewith deny some or all of performing transactions.



Benjamin Marty [permalink](#)

In my experience using the bitcoin client, you are not allowed anything on the bitcoin network until your block chain is in syn latest transactions. It somehow recognizes how far behind you chain is and starts downloading blocks and tells you how old y chain is and how much left you have to update as it download BTW, I un-installed the bitcoin client because over the 1 year s had it installed, the block chain went from about 2 GB to abou and the novelty of having my own copy of the block chain wor comparison to its cost. It would be nice if there were some kin block” that could be generated that flattened the tree into a sir enumerating the value stored at each address.



Michael Nielsen [permalink](#)

On the naysayer DDoS attack on two-phase commit: if someo that they see evidence of double spending, you’d require ther evidence in the form of a signed transaction. The requirement signature makes this hard to forge by a malicious naysayer.




uez [permalink](#)

Here is a very entertaining rational explanation

<http://www.bringhurst.org/2013/04/03/how-does-bitcoin-crypto-work.html>



One thing I still don't fully understand is how the bitcoin reward size is decreased. Who enforces the rules that 25 bitcoins are awarded for validating a block and a few years hence, it'll be 12.5 bitcoins? If we were to decide that the reward should be different (remaining at 25 indefinitely, for example), what exactly would we need to change? Is it the bitcoin mining clients that are hardwired to only validate transactions that award 25 coins to other miners when they validate their block? Or is the date of the validated block indicates that the award should be 25 BTC?



Michael Nielsen [permalink](#)

It's hardcoded, based on the number of blocks in the blockchain. Every 210,000 blocks the rate halves. No need to keep track of the count blocks.




Michal [permalink](#)

Hi,

And cannot miners just continue to validate blocks by adding 25 bitcoins?

As the chain is just a validated list of transactions, is there a cap on transactions?

What does hardcoded mean practically?



Jozef [permalink](#)

You only own the private key (which is hardcoded here, or supposed to be hardcoded) to be able to continue giving the reward to 12.5, I



Nick P [permalink](#)

I hadn't had time to thoroughly delve into the protocol and your excellent work, exactly piece by piece, what/why I needed. The "why's" are extremely important for people who might want to build on top of the protocol as it helps them understand what they should or shouldn't modify.



pebird [permalink](#)

I was thinking about how the blockchain is managed as more transactions are processed, thanks for the link <https://en.bitcoin.it/wiki/Scalability>

Interesting, one of the potential solutions discussed is the use of dedicated nodes instead of lightweight clients to increase transaction rates, reduce latency, and increase blockchain size (via techniques such as "pruning" the chain), etc.

What this implies of course is an evolution into "banks", a group of entities with sufficient resources and staying power to dedicate specialized BTC infrastructure for transaction handling.

In a way, Bitcoin is replicating a history of money evolution in an accelerated manner. I wonder what will take place in the protocol to allow the peer-to-peer nature of Bitcoin while scaling the project to allow the transaction capacity necessary for a global currency.



Michael Nielsen [permalink](#)

Yeah, that is very interesting. I don't run a full client myself, I u
client that doesn't have a full copy of the block chain. And you
see a lot of signs of centralization with the big mining pools:

<https://blockchain.info/pools>



Jeremy [permalink](#)

Great article. Thanks very much for writing it.



Reader [permalink](#)

Typo: requiring = requiring

[MN: Thanks, fixed.]




Reader [permalink](#)

The concept of 'block' (and/or a definition) is not introduced before it is use
makes the concept difficult to grasp.



Reader [permalink](#)

Wouldn't three phase commit be more appropriate than two phase commit



Michael Nielsen [permalink](#)

That would likely be even better, although I haven't thought at
lot of detail.



Reader [permalink](#)

Grammar:

to take account of those asymmetries => to take into account these asymr

[MN: Not sure that was ungrammatical, but in any case I've improved it: "to take into account those asymmetries."]



Reader [permalink](#)

Multiple typos with '...the recipient is receiving 0.39 Bitcoins...' -> 0.319

[MN: As noted earlier, fixed.]



Stuart Quimby [permalink](#)

Thanks for such a generous and informative post. There is so much babble that it often seems to operate socially as more of a rorschach test on current actual means of exchange. I quite agree that the details are considerably more interesting than yet another pundit's babble about what it all means. The details, in my delight, are in the details.

Stuart Quimby



Terikan [permalink](#)

Bitcoin has fascinated me recently. I admit to not being able to fully wrap my head around it, but I took what I could and wrote a little here:

<http://mimictrading.com/viewtopic.php?f=4&t=293>

Maybe you can help me out with one part of this I don't quite get. The sign-off: "Does the block chain know that the address sending the coins is correct?"

sends their sig to go with it, I assume paired up with the hash of the address
the various nodes to validate right?

But if you are sending your sig out then can't any node have access to that
info. They would need to in order to validate. So can a sig only be used once
how is it generated and what prevents it from being faked?



cryptograffre [permalink](#)

https://en.wikipedia.org/wiki/Digital_signature

Public key cryptography is a remarkable and beautiful thing. Even
using Bitcoin has keypairs – one key in each pair is public, the other
private. The nature of asymmetric cryptographic digital signatures
can sign any piece of data using my private key, and anyone with
only my public key can verify that the person who signed that
data has the private key. There's some fascinating mathematics involved
a simple numerical relationship between the public and private keys.



richnormand [permalink](#)

Very nicely done write-up.

Makes me wonder about the news at various times about a major "theft" of
mostly in exchanges. In order to benefit they would have to be converted to cash
introduced later on.

Some of these were for large amounts and not really easy to hide, unless
on them?

Michael Nielsen [permalink](#)



I've wondered the same thing. Some observations: if you copy private key, and then erase their copy, there is no way for them that it was ever truly their key. And if two people both have a copy of a private key, how do you determine who "truly" owns it?

The situation is complicated further by the possibility of laundering. If someone quickly spends some stolen bitcoins, then it becomes very difficult to later recover those bitcoins, since now they may be in the possession of honest parties.



shyam jos [permalink](#)

the best explanation ever , thank you Michael 😊



Austin [permalink](#)

This was a fantastic article and answered all my questions about bitcoins.

Thanks!



someone [permalink](#)

What about the actual code? How many miners are using same piece of software?



Reader [permalink](#)

Indeed, this is a critical question.

The more implementations there are, the stronger Bitcoin would be. It would not be dependent on the "features" or flaws of one particular implementation.

The apparent lack of unambiguous protocol documentation makes me think that alternative implementations are difficult to achieve.

Certainly, it would greatly help if there was some form of “RFC Standard”, or “W3C spec” for Bitcoin.



Besmir [permalink](#)

Hi,

Your article was very interesting and detailed, so I learned a lot more from one question or doubt: What is done with all these hashes? are they gonn for cracking/decrypting encoded data?

what is the real benefit behind generating hash tables?


thank you



muthuveerappan [permalink](#)


Did you do this video or is this video inspired by this post !! ?

<http://youtu.be/Lx9zgZCMqXE> – this too is good...



Michael Nielsen [permalink](#)

I didn't make that. I just watched a few minutes – it looks prett certainly much more detailed and accurate than most of what'



Michael Nielsen [permalink](#)

Many people have asked about scalability, so let me just leave this here:

<https://en.bitcoin.it/wiki/Scalability>

It doesn't address every possible concern, but I think the upshot is that the room for Bitcoin to grow.



Haroun Kola [permalink](#)

Thanks. There's so much to learn about this currency and I'm loving all the that its getting.



Diogo [permalink](#)

Great article! I have a question: Could miners run a modified version of the choose not to publish a transaction in the blockchain? I mean, like a small powerful miners controlling the entire network?



Mark [permalink](#)

If you control half or more of the total mining power in the network can keep a transaction out of the blockchain by solving blocks (average) than the miners who are trying to include that transaction. If you control less than half, you can delay the transaction, but soon the rest of the miners will get ahead of you and your version of the blockchain will lose out.



jshell [permalink](#)

Great explanation — but doesn't solve this problem:

Bitcoins aren't actually backed by anything other than server time.

There was a time in this country when you can go to the bank and trade in a dollar bill for an oz of gold. You can't do that anymore, b/c today dollars are debt not gold.

But bitcoins are backed by server time. That almost makes less sense than debt.



Thomas Mahoney [permalink](#)

I think you're confusing an investment with a medium of exchange. An investment should be "backed up" by something, in the sense that it should give the holder a claim on future cash flows or other real assets. A medium of exchange is just that, something used to facilitate transactions. It's an accounting device. It should have scarcity value and be resistant to counterfeiting. Fiat currencies have scarcity value to the extent that they are usually printed in finite amounts. Gold is generally scarce. Bitcoin is scarce as well.

Gold has been used as a medium of exchange for centuries. What's "backed up" by? Nothing. It's just scarce, and therefore suitable as a medium of exchange.



Miksa [permalink](#)

BBC had an interesting article called "Value" 2 years ago.

<http://www.bbc.co.uk/news/magazine-20120228>



Marco [permalink](#)

If people are willing to pay for something that is rare or unique, it has value.


A Ford Mustang '65, first issue Marvel Comic books, baseball cards, Bitcoins all have value because they're scarce and people are willing to pay for them. The demand for it defines the price.

Marc [permalink](#)



Excellent write-up, and I look forward to further installments – which leads are you no longer updating your RSS feed(s)? I came here from Bruce Sc blog, and I like what I see so I subscribed (in [goread.io](#)), but neither of you anything newer than the beginning of 2013.

I'd like to clarify: I'm grateful for your posts, and I'm not complaining if you' the whole RSS thing (Google did, why shouldn't you?). But if you _haven't RSS, but it was supposed to be getting updated automatically... it isn't.



Michael Nielsen [permalink](#)

Thanks for pointing this out.

I just checked both RSS feeds, and they seem to be fine. I typ longer essays, often in the 3,000-20,000 word range, which is update my blogs a few times a year. You may enjoy looking th of my past articles. This blog carries my more technical stuff, \ other blog (<http://michaelnielsen.org/blog>) is more general.

Your comment did make me notice and fix some mistaken link sidebar, so thanks for that!



Marc [permalink](#)


Huh. I clicked through to the Feedburne there. Perhaps the problem is on the go so far, but... Maybe later I'll try again w

Here's what I see for the DDI feed: [http](#) and here's what I see for your main fee

I've tried refreshing multiple times, but

I've bookmarked your pages and I can periodically to see if you've got anythin

to know that (at least for some of us) R



Michael Nielsen


I checked in my
come through fir



Joseph [permalink](#)

Am I to understand that it takes about 60 minutes to pay somebody through the Bitcoin network? I reached this conclusion based on the 10 minute average confirmation and the requirement of it being 6 back in the chain before it is confirmed.

Thanks,
-Joseph



Michael Nielsen [permalink](#)

Full confirmation requires about 60 minutes. Many people are willing to accept payment on more trust, though, say after just a single confirmation (~10 mins).

Yes, this is a significant disadvantage of the protocol in its pre-confirmation architecture.



Marco [permalink](#)

Transactions are instant. Confirmations are not.

A confirmation takes 10 minutes. If you want full confirmation, on average it takes an hour (6 confirmations).

For eCommerce, this will probably work in most cases.

For retail, this can be an issue.


However, there are a few points:

- If you try to double spend, it doesn't mean you will succeed.
- It is not easy to double spend in front of the cash register (you have to build some app and be all prepared)
- The merchant (I think BitPay is doing this) can listen on the network to see if there was a double spend attempt. Those are easily detected.
- Don't forget that a Credit Card payment can be charged back later. Just saying. 😊



Gregory Johnson [permalink](#)

Love the article. It is the first article that I have been able to understand on Bitcoin and I have been reading a few on it. And as a comment to style, I really appreciate a higher-principled discussion on the topic. I am so annoyed with the internet barrage of get-rich-quick articles on this, or the excitement of the exchange. I have yet to read before now any intelligent comments to the social value, in my opinion. Your link to <http://szabo.best.vwh.net/formalize.html> was appreciated.



Michael Nielsen [permalink](#)


Thanks. All of Szabo's writing is worth reading, incidentally, his is a treasure trove.



Steffen [permalink](#)

You know that there are some hints that
Nakamoto?

<https://likeinamirror.wordpress.com/2011/01/14/steffen-probably-nick-szabo/>



Bart [permalink](#)

great article.



Ian [permalink](#)

Re: why BT doesn't use 2PC, as I understand it, it's because 2PC becomes exponentially more complex/unreliable with an increased number of parties.

2PC is a collapsed version of the byzantine / paxos protocols (which is 2P place of 2), and the basic problem is that a lot of nodes have to be online & a lot of messages in order for it to be workable. In any case it's more complex than 'longest blockchain wins'.

That's my amateur understanding anyway.




HPublius [permalink](#)

Great article and great discussion! This is a very good overview of the technical aspects around the bitcoin protocol. The fact remains that bitcoins have no intrinsic value and the promise of a peer-to-peer payment network (medium of exchange) cannot be fulfilled unless the bitcoin is transformed into a true digital currency. Here are my thoughts on how to accomplish that: <http://tinyurl.com/m57hd2z>



OnlyMe [permalink](#)

Hi,
first of all great explanation on Bitcoin, I love it!
I guess my question is simple to answer.
How can I verify that a transaction is signed by a certain address if all I got
of the public key? Don't I need the full public key for that instead of only th
What am I missing guys??




Michael Nielsen [permalink](#)

The transaction contains the Bitcoin address of the payee (or
there are multiple outputs) in the output fields, and the public k
signature(s) of the payer(s) in the input fields. So there's no pr
you do have the full public key of the payer.



kgb [permalink](#)

Silly question from a non technical person: how will transactions be approv
verified subsequent to 2140 when there are no more rewards for mining?



Michael Nielsen [permalink](#)

Transaction fees (which I briefly describe in a couple of places
article).



Tom Hatcher [permalink](#)

Thank you. Best explanation I've seen so far! I still don't understand it completely but it's slowly becoming clearer. One question, though. I hear that it's open so we can look at the source code. I'd like to do that. It's written in C++? Where is the code and look at it? Thank you.



Tom Hatcher [permalink](#)

I answer my own question. It's at github.com



Fee Fi Fo Fum [permalink](#)

I would love to see you discuss tumblers and the effectiveness and possibilities of anonymizing your bitcoins




Cb [permalink](#)

In your anonymous section you speak of debunking a fairly huge myth with backing it up. You just state the equivalent of "actually it's not anonymous" going into detail.

I don't believe that to be true unless you are implying that various ways of internet anonymously are breakable.

For instance if TOR is compromised versus if it is not, or if other methods of traffic surrounding use of bitcoins are insufficient

Your assertion that bitcoin is open and transparent has nothing to do with its use it anonymously, and the claim that it wouldn't be able to 'achieve convincing' will narrow the pool of suspects down sound closer to a statement of successful anonymity rather than unsuccessful.



Michael Nielsen [permalink](#)

“You just state the equivalent of “actually it’s not anonymous” without going into detail.”

It’s certainly not meant to be a proof! I do, however, go a great deal further than just saying “it’s not anonymous” — I reference a large and growing body of academic literature that takes supposedly anonymous transactions and then de-anonymizes them. I believe techniques similar to those in those papers will be very useful for attacking Bitcoin. There are many complications in Bitcoin, notably that some people (though far from all) routinely use new addresses for each transaction. That makes de-anonymization an interesting challenge, and (I think) is different than in earlier work on anonymization. I’ll be most curious to hear what the de-anonymizers have to say after making a sustained attempt at Bitcoin.



Magnus Sorenson [permalink](#)

Linking bitcoin addresses to a real identity is somehow associated with an address in the street with cash – without revealing the public network, and a device with a name, then I can spend that bitcoin with a personal info to whoever I send that bitcoin. The guy who sends me the drugs would not be fudged as well. If he does not send the drugs, I get linked to me if the drugs arrive safely on that transaction, not even the NSA could do that. Now if I do the same thing many times, I can be linked to me through other vectors – but explain how anyone can link me if I buy the bitcoin with cash on the street.

permanent record. I assert that bitcoin forever- if you do it right.



MerkleTree [permalink](#)

The system is all somewhere and track addresses during man trans There are sever

I suggest you lo

You will find that effort.


To remain anony This includes the they will work.

This leads to the dangerous to sp



Adam Back [permalink](#)

You mention using multiple sub-puzzles to reduce variance. This is a bad i introduces progress. Because bitcoin is a first-past-the-post race, progress unfair advantage to more powerful miners (>2x reward for 2x power).



Michael Nielsen [permalink](#)

You're presuming a particular design (in which people solve a puzzles privately). There is no necessity to make that presum

why I make the first part of the statement “with some careful design it is possible to considerably reduce the variance in the time to validate a block of transactions”. I haven’t worked out the best possible design yet, however, I have found a design whose distribution is quite a bit more sharply peaked around 10 minutes than the current puzzle. Unfortunately the details are more complex than I want to write out right now, but I will come back to it in a future post.



Adam Back [permalink](#)

Also you talk about risk of nonce reuse. This won't happen because people use their own reward address, so even if the nonce is reused the work proof will be invalid.

Further in the case of pool mining the pools hand out work, specifically to a range of nonces (which is somewhat insecure as others could guess the work range and race them to produce it).

And finally the secure way is pooled miners use `getblocktemplate` and use a random counter start extranonce. If extranonce is large enough and random enough the probability of nonce collision is practically 0. You can read about this in the paper <http://hashcash.org/papers/hashcash.pdf> or <https://en.bitcoin.it/wiki/Extranonce>

For decentralization miners should also choose their own blocks by running a node and filling in the details into the coinbase provided by `getblocktemplate`.



Adam Back [permalink](#)

two-phase commit: if you are willing to wait 10 minutes, bitcoin already does a two-phase commit. I presume the form it would take is the proof of double spend would be (or rather, the proof of no double spends).

There have been proposals to forward double-spends with a double-spent transaction (currently the first only is received).

Maybe just an api to ask if there are any transactions conflicting with a given transaction a user could ask a few random nodes to gain confidence.


You also have to bear in mind preserving the 0-confirmation spend function people rely on that for low value point of sale transactions.



SRM [permalink](#)

“I’ll return later to the question of why the Bitcoin address is a hash, and not a public key.”

I don’t think you ever come back to this topic.



Michael Nielsen [permalink](#)

See the discussion above, in reply to Benoit Mason.



JohnT [permalink](#)

Thank you for the primer.

You might consider removing the footnote. IMO, Bitcoin cannot be successfully defended as free speech. Free speech is not a full blown unlimited right, and in a crowded theater reminds us.

But Article I, Section 8, subparagraph 5 does grant Congress full power to “coin money and declare [its] value.” And a subsequent subparagraph grants Congress power to outlaw any currency it wishes for citizen uses as legal tender.

Thus, IMO, the Supreme Court could never allow free speech to prevail over Congress’s unfettered Constitutional authority.

I hope this comment does not derail a great discussion of Bitcoin. Please delete my comment if it becomes a red herring.



johnp [permalink](#)

a question

re the part about

“”

everyone (collectively) is the bank. In particular, we'll assume that everyone Infocoin keeps a complete record of which infocoins belong to which person

“”


assuming millions and perhaps eventually billions are using bitcoins

so for example billions of transactions could take place daily

times this by the qty of bitcoins (each single one being unique)

times this by the billions using bitcoins

then what affect would that have on the network ?




Michael Nielsen [permalink](#)

Already addressed in comments above.



Pavel Masyuk [permalink](#)

That question about a nonce... I think that the parameters of the puzzle change every single miner. Everyone's desired block contains a unique transaction that no other miner has – a transaction of giving a reward to himself. So there is no trying to trick others – parameters of their puzzles are different. It's OK for you to just try 0,1,2 etc...



Michael Nielsen [permalink](#)

Already addressed in comments above.



Joshua Holden [permalink](#)

So if I've got this right, one proof-of-work computation takes about 10 minutes. You (currently) get 25 Bitcoins for doing it, and each Bitcoin is (currently) about 1000 USD. Right? If so, this only makes sense if most proof-of-work computations don't get finished and/or don't get rewarded. Is that usually the case? Does someone else get there first? Do you know about what fraction of proof-of-work computations get rewarded?

This was really useful; thanks!



Steffen [permalink](#)

Joshua,

yes. It's a race. Whoever finds the hash that is smaller than the target (defined difficulty), they will gain the reward for the block. The difficulty is adapted every two weeks or so to reflect the changing (now growing) power of the network.

The power is growing so fast and so much that some already call it an "arms race".

-st



Chris Crozier [permalink](#)


Thanks for the article, very interesting, but I don't see any mechanism for combining fractional bitcoins.

I'm puzzled by what seems to be an ever-increasing fragmentation. It seems like over time you would accumulate a large number of coins of varying fractional values. To make a payment you would have to lump together a collection of fractional coins until they equal or exceed the transaction required, then typically end up with paying

your change. This one-way process of cutting off pieces of a bitcoin would steadily. A holding of one bitcoin would end up being constituted of maybe thousands of differently-sized fractions.

If that's the case, that will make transactions increasingly messy: you may consolidate a large number of inputs for one payment. In turn, that will lead block chain file growing faster and faster.

Did I miss something?



Michael Nielsen [permalink](#)

This is not a problem. Yes, some transactions will “fragment” but other transactions undo fragmentation. For example, a 5-input transaction will reduce fragmentation.

This sounds a little complex for the user, but in practice, good software will make this invisible. You simply say “I want to send such bitcoins to so-and-so address”, and all the details of combining transactions will be taken care of. In this sense it's actually easier than cash, where we deal with the fragmentation / de-fragmentation at the time (i.e., finding the right combination of bills and coins to make service, and then dealing with the resulting change).



Amit Prakash [permalink](#)

Amazing article. Looking forward for more on similar topics – may be someone can explain for me.



Edward [permalink](#)

Very informative article; thank you. I was wondering:

1) With regard to transaction fees, I assume it is up to the first miner to successfully validate a block to determine if your offered fee is large enough to be included in the next iteration of the block chain. Additionally, I assume you have to determine if you're willing to pay well in advance of the next chain being verified. If this is the case and we fast-forward to 2140, won't all miners (assuming the computing power is concentrated) be incentivized to take ANY transaction fee no matter how small? For instance, if I have .001% of the computing power I should, on average, validate a block once every 2 years (hopefully the math's right, it would be infrequent in that regard) and, considering (assuming?) there would be no marginal cost to include a transaction (or 1,000 transactions) offering me only \$.1 equivalent in transaction fees, I can't think of why I wouldn't validate that transaction. As such, even if there are a few big players who controlled say 2/3 of the computing power, on average, blocks would be validated by a smaller player who wouldn't care about price. If people already wait nearly an hour to finalize a transaction likely they'd be willing to wait an extra ~30 minutes and, as such, the big miners would likely just lower their price thresholds pretty substantially; Wouldn't this create an odd prisoners' dilemma situation? Maybe I'm missing something but it strikes me that this would encourage fees forcing people out of business until computing power concentrated among an oligo(mono)poly of miners who could exercise sufficient pricing control which, as you pointed out, would probably create integrity concerns.

2) I believe I've previously read somewhere that there was a price threshold to incentivize mining at different levels of complexity – for instance, if bitcoin price fell tomorrow, would miners continue to mine or would the marginal cost of running old equipment outweigh the reward (ignoring fees)? If that's the case and a scenario where that occurs, does bitcoin grind to a halt or will some miners shut down or switch to less expensive equipment as they did when bitcoin was in that price range? This sort of thing boils down to whether the use of high cost computing equipment is a function of competition (and price) or problem complexity?

Both are hypothetical but I was curious to know if you (or anyone) had considered these questions.

Thanks again for the article.



Edward [permalink](#)

(I think) I figured out the answer to #2 – I was unaware of how was calculated.



JimmyWeg [permalink](#)

Great article! I'm working on a case and see that the bitcoin user employs a dozen different applications: Anoncoin, Phenixcoin, Primecoin, etc. I take protocol is the same among the clients, though hash algorithms, proofs of and the like may differ. From what I understand, if I use XPMs and want to something from a vendor who accepts BTCs, I have to go through some b exchange facility to complete the transaction. If that's correct, and consider U.S. in a vacuum, isn't it like we're all carrying around a different brand of (USDs, Yen, Pounds, etc.) and have to exchange them almost every time buy something? It's easy with credit cards, but I don't see a similar approach bitcoins. Thanks!



David [permalink](#)

"I don't understand why double spending can't be prevented in a simpler n using two-phase commit. Suppose Alice tries to double spend an infocoin Bob and Charlie. The idea is that Bob and Charlie would each broadcast t respective messages to the Infocoin network, along with a request: "Shoul this?" They'd then wait some period – perhaps ten minutes – to hear any r who could prove that Alice was trying to double spend. If no such nays are provided there are no signs of attempts to disrupt the network), they'd ther transaction. This protocol needs to be hardened against network attacks, t to me to be the core of a good alternate idea. How well does this work? W drawbacks and advantages does it have compared to the full Bitcoin proto

I think something along these lines is planned:

“On top of all that is a long list of new features and improvements I’d like to put into a 0.9 release; the highest priorities on my wish list are:

1. “First double-spend” relay and detection. Detecting attempted double-spending as soon as possible is great for low-value, in-person transactions, and we should work more to support that use case.”

<https://bitcoinfoundation.org/blog/?p=204>



David [permalink](#)

Great article. I will use it as a self-study tutorial.

In your next instalment, could you give a broad description of where the protocol is actually to be found (is it a particular piece of software?), how can it be changed, who can change it and indeed to what extent is it capable of being changed? These are important questions because they go to the ability of Bitcoin to evolve and adapt, but it is very hard to find any good general account of these issues.

Best regards

David



Nitán Shalon [permalink](#)

If Satoshi Nakamoto already made two patches to Bitcoin, what’s stopping him from making another patch right now that destroys Bitcoin?



Arthur Colle [permalink](#)

“Of course, after Alice has published her message it’s possible for other people to duplicate the message, so in that sense forgery is possible. ”

How can someone forge the message without Alice’s private cryptographic key?



Dave [permalink](#)


Hi. Great article. Thanks for putting it together.

I am still having one big problem — and I feel like I must be missing something obvious. You wrote:

“Suppose Alice tries to double spend with Bob and Charlie. One possible scenario is for her to try to validate a block that includes both transactions. Assuming 10% percent of the computing power, she will occasionally get lucky and validate a block by solving the proof-of-work. Unfortunately for Alice, the double spending is immediately spotted by other people in the Infocoin network and rejected, forcing her to solve the proof-of-work problem. So that’s not something we need to worry about.”

Who is going to be looking to reject it, and what does that even mean? If a miner (Alice) manages to complete a block that contains transactions that are in fact, valid then what? Do other miners check them before building on top of the block? And, if not, then what does it mean for others to ‘spot’ them.

This has been bugging me for days! Your thoughts would be greatly appreciated.



Michael Nielsen [permalink](#)

Anyone with a copy of the block chain is not going to accept a block which has an obvious attempt to double spend in it.




Dave [permalink](#)

Thanks!

So that means that miners examine each block and choose to build on it...? And, if that’s true, then the inspection built into the incentive structure means that a block becomes invalidated if down the line someone finds a double spend.

a block with a double spend? (That would be the first one that I have seen described.)



Michael Nielsen

Checking for a c



Daniel Park [permalink](#)

That's a long detailed process of bitcoin above which was a great read. I don't understand near the end, but I definitely get the gist of it.

I don't perfectly understand the relationship between transactions and blocks (are blocks only required when ppl try to cheat the system?) If so, the money earned by miners is essentially imaginary and something that only exists within trust that bitcoin will continue to work. I guess there are two cases:

1. as miners get bitcoins, the value of bitcoin compared to other currencies decreases (This is probably the case)

OR

2 as miners get bitcoins, if it comes to a point where everyone wants to convert bitcoins to real currency, it's not going to equal each other; bitcoins -> convert to real money at a bitcoin exchange rate -> not enough real money.

Also, why assume every 210,000 blocks is occurs every 4 years? is this an assumption based on bitcoin flow so far? Wouldn't every 210,000 blocks occur more often if there is more flow?



Steffen [permalink](#)

If everybody would like to exit Bitcoin at the same time the price would collapse. The current speculation is though, that the opposite

Many people try to buy bitcoins for the fiat money. How many? here <http://fiatleak.com>.

With regards to why 210,000 blocks are created in roughly four years. The network difficulty is set so that only six blocks per hour can be created. Roughly every 10 minutes a new block enters the blockchain.

Let's do the math: 4 years \rightarrow 1461 days \times 24 hours \rightarrow 35,064 hours \rightarrow 210,384 blocks




Jozef [permalink](#)

Fairly good explanation although some important things missing. One mistake is that the difficulty is not constant.

“So if we want the output hash value to begin with 10 zeroes, say, then the difficulty is 16¹⁰. On average, you need to try 16¹⁰ (approx 10¹²) different values for x before you find a suitable nonce.”

I think it should be 16⁹ on average. At worst it is 16¹⁰.



Michael Nielsen [permalink](#)

No, the average number of trials is 16¹⁰.

(If we repeatedly sample from a Bernoulli random variable with probability p of a success, then the expected number of trials to success is $\sum_{n=1}^{\infty} n p (1-p)^{n-1}$, which is easily shown to be $1/p$. In the proof-of-work is 16⁻¹⁰, so the expected number of trials to success is 16¹⁰.)

What important things are missing from the explanation?

Lex [permalink](#)



i have alot of questions about bitcoin and ive done alot of searching, this e has to be the one that has got me the closest to the answers ive been look unfortunately im still falling short of understanding alot of the basic, i do wa question here which seems to be the most important question for me. In co science terms what is a “bitcoin” there were points in this article where it s close to specifying this, although i may have actually missed the answer, s “bitcoin”?? is it a unique value? if i have 100 bitcoins what is it that i really unique values ? or is it simply a value that was “said” to be giving to me ar assigned a value, sort of like the 5th entry in a ledger, in the genesis block “said” 50 bitcoin was givin to whom ever,50 bitcoin was given to bill and w that 50 1234567.. ??? overall im still searching for clarity on the fundamen bitcoin from a computer science view



david lee [permalink](#)

Michael,

We are organising World Bitcoin Conference on 24 & 25 March 2014 in M
Would like to invite you to speak on Bitcoin protocol. Could we have your e
address to send the invitation.



David [permalink](#)

Hi – what a great write-up! As for your first Author’s question – There are p
several less complex methods for confirmation but there is inherent securi
current approach which appears to organically solve the problem without r
on factors and layers outside the network itself.



momingle [permalink](#)

There is one part that I am not sure I understand. For example, let say mir
his queue transaction A, B, and C to validated on a new block. Is it possibl
will have transaction B, C and D in his queue (but not A) that he will valida

new block? Assuming both solve the puzzle. Now both transactions B and two different blocks. Will both blocks got accepted?



Nagaraj Hubli [permalink](#)

Thanks for the write-up, it helped me a lot in understanding the underlying Bitcoin protocol. Can't wait for the next in series.



elvin [permalink](#)

The protocol rules in the bitcoin wiki are ambiguous when an incoming block designates as its predecessor a block somewhere on the main branch, what happens exactly ?

the wiki can be interpreted in two ways :

1. nothing at all – the incoming block is NOT added as the beginning of a new branch
 2. the incoming block IS added as the possible beginning of a side branch pending any verification for the moment. If that is the case, then most blocks in the branch should have many “brothers”
-



elvin [permalink](#)

I have verified that the correct answer is 2.

I have also found what limits the number of “brothers” for a block:

1. miners stop mining as soon as they receive a valid node and no new blocks are sent
2. their neighbours do not relay the blocks they might have seen during the meantime

Sorry for asking

fivebuckchuck [permalink](#)



Thanks for the write up ..

How do transactions get organized into blocks ? Are nodes broadcasting transactions or blocks ?

Number of transactions in a block ; is that hardcoded in the protocol ?

in your example of $h(\text{"hello world"}|\text{nonce})$, is “hello world” a unique transaction or the whole block ?



Rich [permalink](#)

“Suppose Bitcoin mining software always explored nonces starting with $x = 1, x = 2, \dots$. If this is done by all (or even just a substantial fraction) of Bitcoin miners then it creates a vulnerability. Namely, it’s possible for someone to improve their odds of solving the proof-of-work merely by starting with some other (larger) nonce. More generally, it may be possible for attackers to exploit any systematic patterns in the way miners explore the space of nonces....”

This is incorrect:

Because the block hash is dependant on the contents of the block. For the possibility of a miner improving his odds through this method the miner must be mining the exact same block as someone else including not using his own coinbase and transaction fees to go to. Removing the entire incentive for

Trackbacks and Pingbacks

1. [How the Bitcoin protocol actually works | Enjoying The Moment](#)
2. [Understanding the Bitcoin protocol | Minsheng 民生](#)
3. [Science & Technology Roundup: #3 | Great Wall of Numbers](#)
4. [mid-day-ish links dump | imagine art -||+ there about?mid-day-ish links dump](#)
5. [How the Bitcoin protocol actually works | DDI | geistwc](#)
6. [How the Bitcoin protocol actually works | DDI | Mark Solock Blog](#)
7. [How the Bitcoin protocol actually works « adafruit industries blog](#)
8. [The Bitcoin Protocol - iamronen](#)

9. [Links 12/8/13 « naked capitalism](#)
10. [Bitcoin for dummies – Author walks users through how Bitcoin actually works | My Blog](#)
11. [Want to Spend Bitcoin in France? \(Caveat Emptor\) | FrenchNewsOnline](#)
12. [Robert McGhee » December 8th](#)
13. [¿Cómo funciona realmente el protocolo del Bitcoin? \[ENG\]](#)
14. [BitCoins](#)
15. [Bitcoinizzle For Shizzle | BubbleCoin](#)
16. [How the Bitcoin protocol actually works | Blog of Leonid Mamchenkov](#)
17. [Starting of “Gift Economy \(reputation based\)” | l'ambizione dell'artigiano](#)
18. [Headline News from PaymentsNews.com – December 9, 2013 - ResidualMarketplace.Com](#)
[Residual Listings for Free](#)
19. [How the Bitcoin protocol actually works « The Bitcoin Channel](#)
20. [Bitcoin Explanation « Electronic Padlock](#)
21. [Bitcoin Explanation](#)
22. [Bitcoin | Thoughts on computing](#)
23. [Links: Mandela and Great Leaders; WTO deal in Bali; How Money Works](#)
24. [How Does Bitcoin Work? ← Clinton Pavlovic](#)
25. [Following Bitcoin Believers is still an act of faith | Digital Trends](#)
26. [Virtual Mining Bitcoin News » Bitcoin Believers aren't wrong, but don't mistake their faith for](#)
27. [Open thread: Bitcoin and similar currencies](#)
28. [How Bitcoin Actually Works](#)
29. [Bitcoin Speculations | Aleph](#)
30. [How the Bitcoin protocol actually works | The Freedom Watch](#)
31. [How the Bitcoin protocol actually works | DDI - The Ripe Group](#)
32. [Explaining Bitcoins | kef's blog](#)
33. [A Detailed Explanation of the Bitcoin Protocol – Glimpse From a Height](#)
34. [Friday Links | Meta Rabbit](#)
35. [Bitcoin](#)
36. [Virtual Mining Bitcoin News » Bitcoin: Open source money whose time has come](#)
37. [Bitcoins, Cash and Gold | GusBook](#)
38. [Plaidoyer pour le Bitcoin | ComplexitésComplexités](#)
39. [Occupy BitCoin](#)
40. [How to not get rich from bitcoin | Turing's Invisible Hand](#)

41. [Bitcoin: The Economic Frontier of a Digital Age | The American Statesmen](#)
42. [Why You Need to Own Bitcoin | Winning Bitcoin](#)
43. [Bitcoin: a bit risky | Learn About Finance Careers](#)
44. [Various well written articles/sites | GSR8 Blog](#)
45. [Bitcoin transactions explained; How the Commodore 128 was ... | BitCoin Razor](#)
46. [Digitopoly | Time for a little Bitcoin discussion](#)
47. [How and why BitCoin will plummet in price](#)
48. [How the Bitcoin protocol actually works and what that means for us | Helping Public Market](#)
49. [| Explaining bitcoin in a bar](#)
50. [Cryptocurrency Cat-and-Mouse games in China | Great Wall of Numbers](#)
51. [Random thoughts: Bitcoin](#)
52. [Info On Bitcoin | Curly's Blog](#)
53. [My Favourite Bitcoin Bookmarks | Webs Wonder Design - Webs Wonder Design](#)
54. [Facecoin In More Detail | Rob Myers](#)
55. [Fed Banker Tries Criticizing Bitcoin, Ends Up Perfectly Describing The Fed – Bitcoin Magaz](#)
56. [Bitcoin 2.0: Distributed Corporations, Derivatives, and Information Markets | The Ümlaut](#)
57. [The Difference between Bitcoin and OpenTransact | OSCurrency](#)
58. [j.r.mchale : The Nuts-and-Bolts of the Bitcoin Protocol](#)
59. [Fed Banker Tries Criticizing Bitcoin, Ends Up Perfectly Describing The Fed](#)
60. [Sécurité Bitcoin #1 : Introduction | Le Comptoir Sécu](#)
61. [De week in 375 woorden « Bits of Freedom](#)
62. [Bitcoins the hard way: Using the raw Bitcoin protocol | Bitcoin](#)
63. [On Bitcoin | Emergent Chaos](#)
64. [bitcoin critical mass/explosion, wheres it going? | Turing Machine](#)
65. [Bitcoin Beyond Money | SoshiTech – Social Media Technology – Soshitech.com](#)
66. [Bitcoin reading for developers | Coffee Spoons of Code](#)
67. [The Homebrew Litecoin Mining Project | Record important events of digital currencies](#)
68. [Complete Dogecoin Guide - Introduction into Mining - Everything Doge!](#)
69. [How the Bitcoin protocol actually works | DDI | Arca](#)
70. [bitcoin | jamesmath](#)
71. [Fun economic reading | MuddyHorse Farm and Tech](#)
72. [Bitcoin: have some of us been had? | Jim Bennett](#)
73. [“The Bitcoin Protocol” Infographic |](#)

- 74. [I don't get Bitcoin... but I want to... | StartupCafe](#)
- 75. [Daily BTC – PBOC edition : Token Intelligence](#)
- 76. [Enter The Blockchain: How Bitcoin Can Turn The Cloud Inside Out | TechCrunch](#)
- 77. [Quora](#)
- 78. [Bitcoin Family Tree | Randolog](#)

Comments are closed.